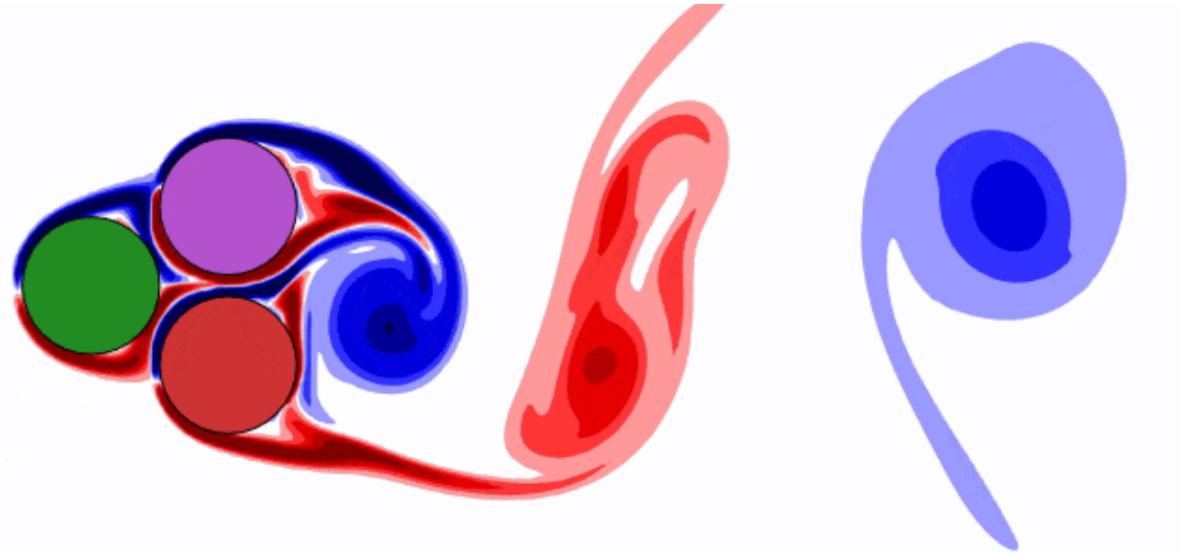


WaterLily.jl: A differentiable solver to simulate fluid flow and dynamic bodies with heterogeneous execution

Bernat Font*, Gabriel D. Weymouth

JuliaCon 2024

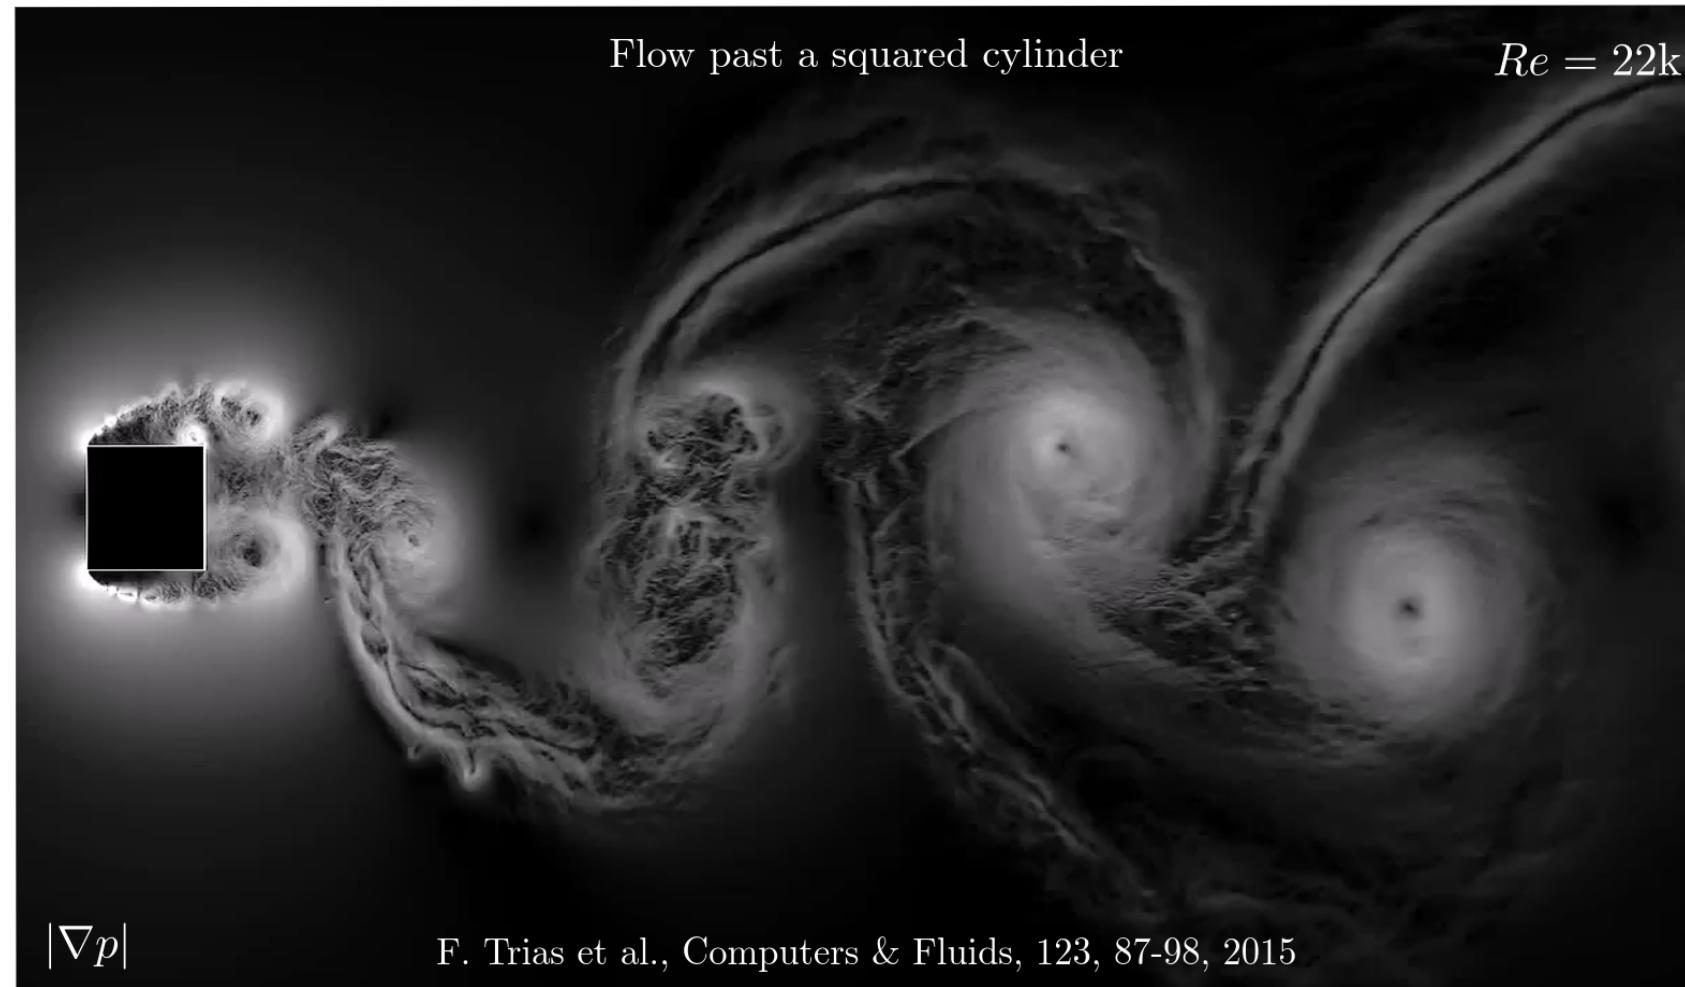


Background | Simulating fluid flows

- Solving the Navier-Stokes equations, aka. Computational Fluid Dynamics (CFD)

$$\frac{\partial \mathbf{u}}{\partial t} + \mathbf{u} \cdot \nabla \mathbf{u} = -\nabla p + Re^{-1} \nabla^2 \mathbf{u}$$

- ▶ Designing vehicles (aircrafts, trains, cars, boats...), weather forecasting, combustion, even biological flows
- But it's costly! There is a wide range of spatial and temporal scales that need to be resolved → HPC



Julia for CFD?

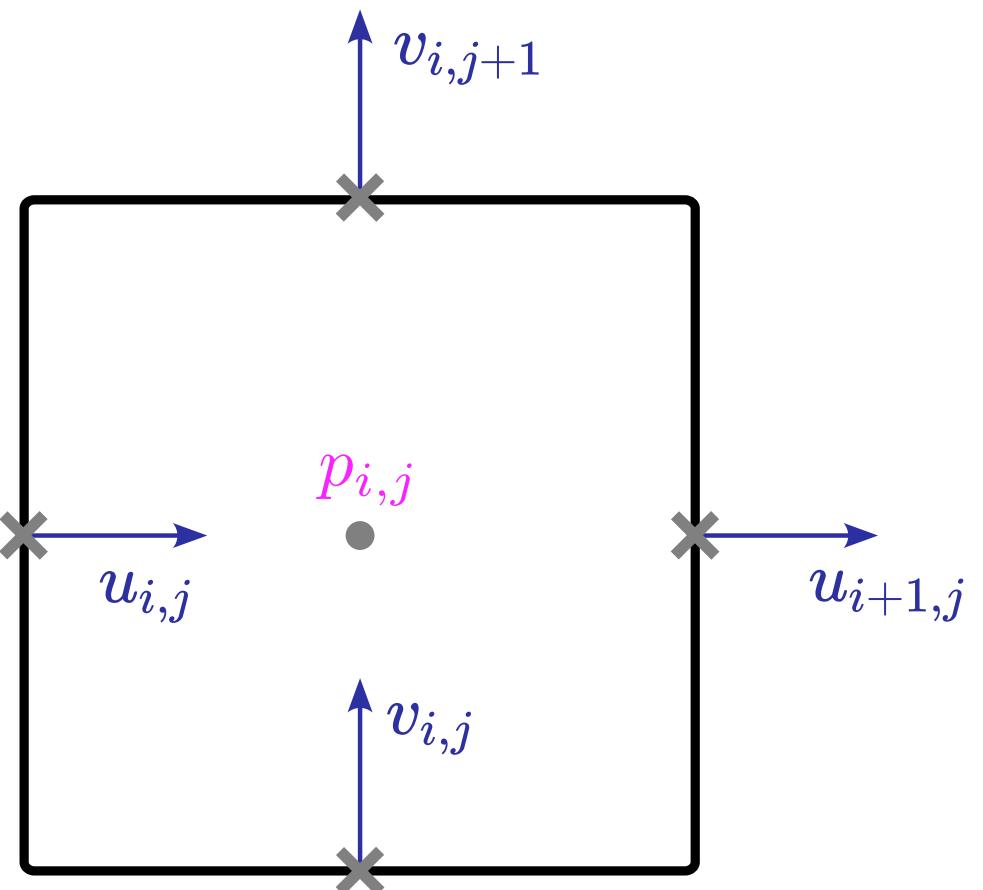
- CFD solvers must be written in compiled languages.
- In CFD, HPC is a must.
 - ▶ Julia is parallel: `LoopVectorization.jl`, `Polyester.jl`, `MPI.jl`, ...
- Modern CFD solvers should be able to run in both CPUs and GPUs.
 - ▶ Julia can run in most architectures: `CUDA.jl`, `AMDGPU.jl`, `KernelAbstractions.jl`, ...

Added benefits:

- Solves the two-language problem.
 - ▶ Easy to integrate with ML libraries such as `SciML`, `Flux.jl`, ...
- Active community.
 - ▶ Engage with core developers and benefit from latest developments

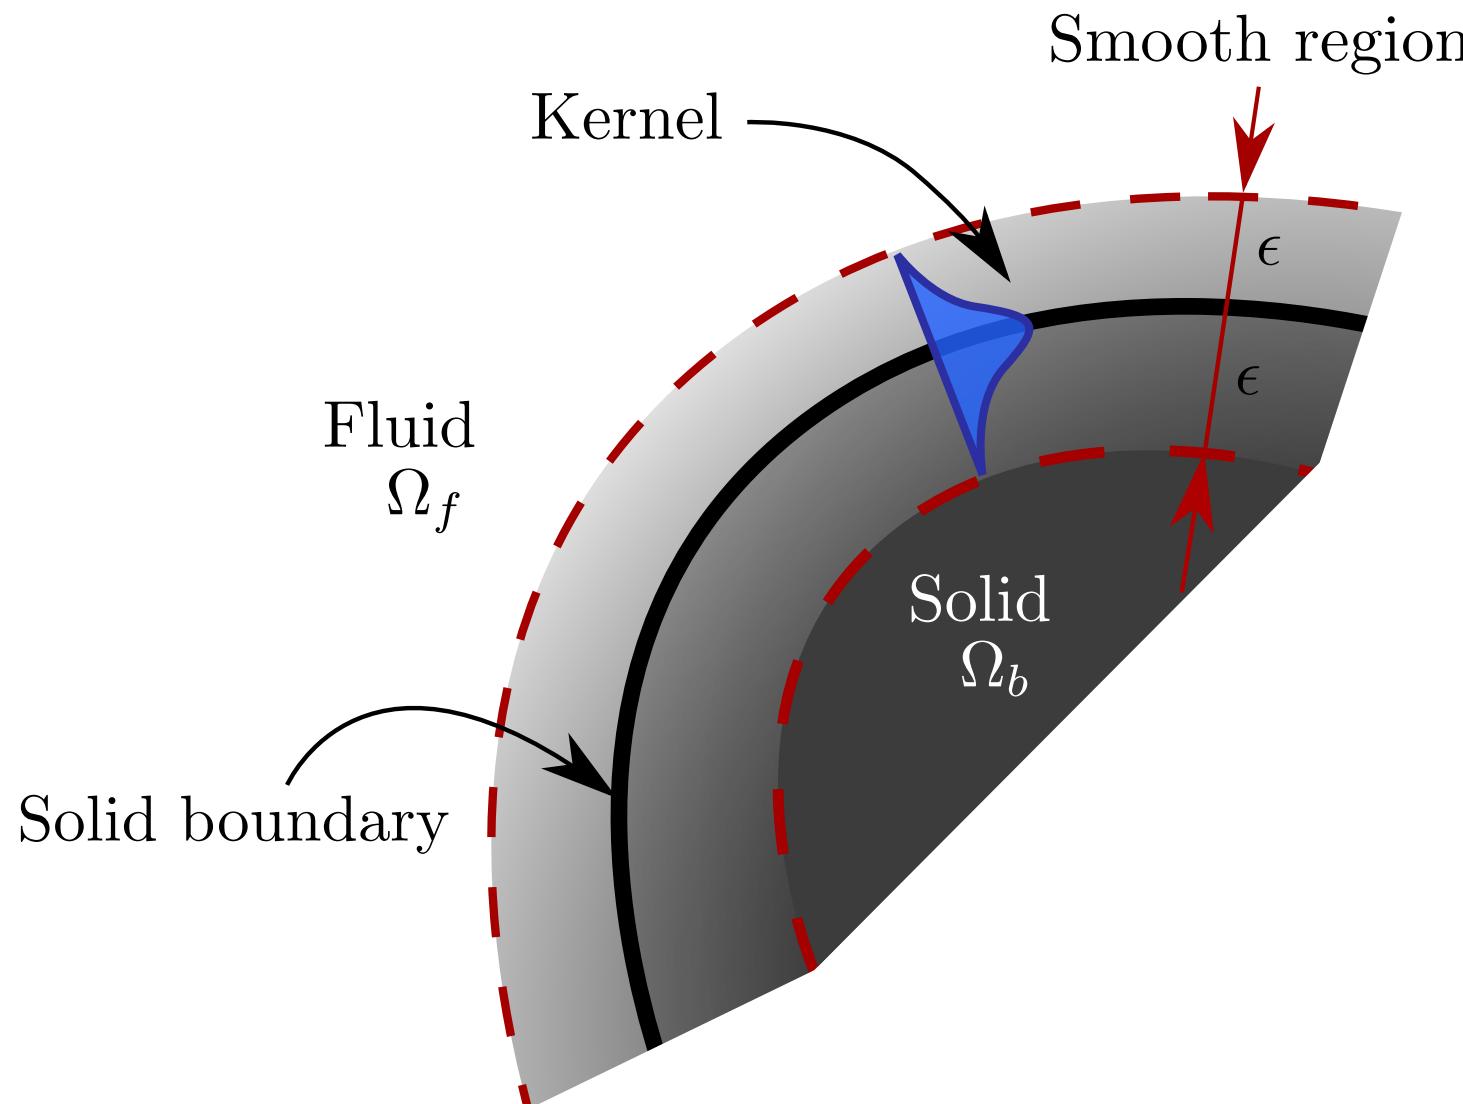
Numerical methods in WaterLily.jl

- Incompressible fluid flow
- Finite volume in a staggered grid



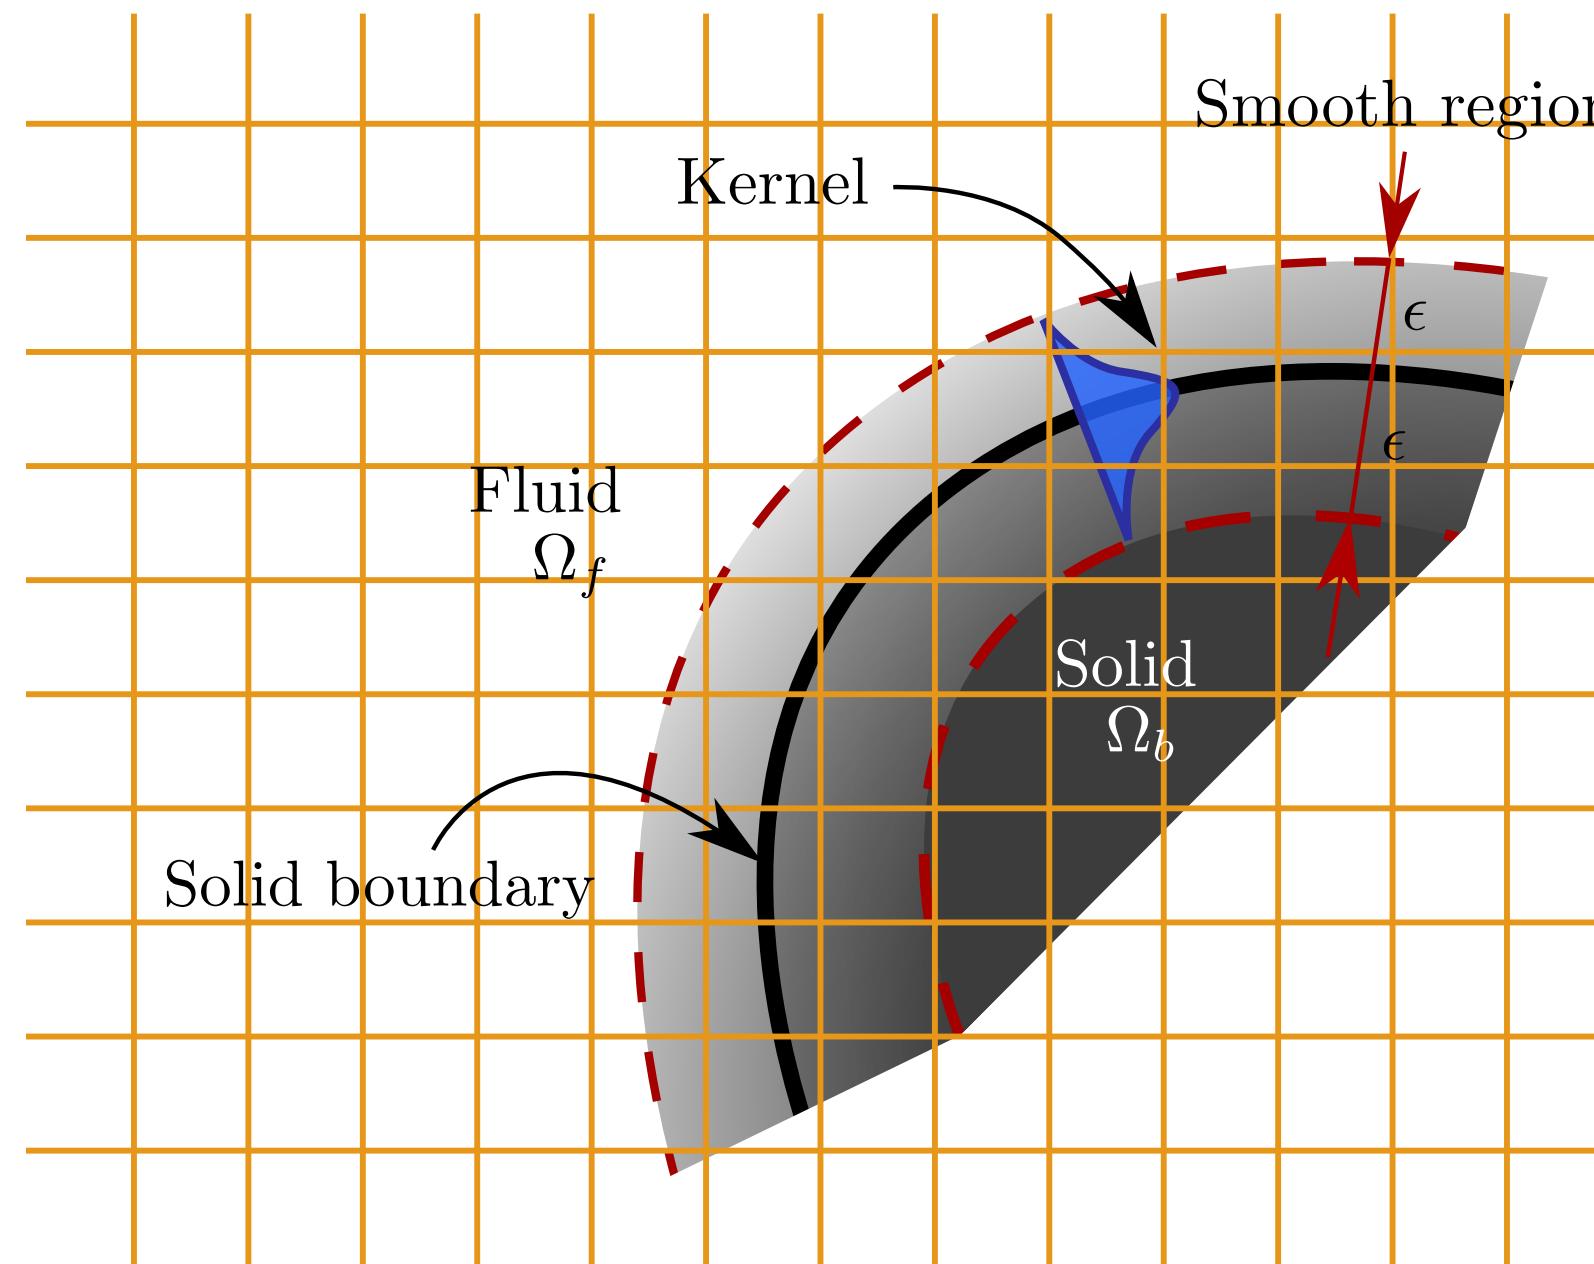
Numerical methods in WaterLily.jl

- Incompressible fluid flow
- Finite volume in a staggered grid
- Boundary data immersion method for solid geometries
 - Signed distance function (SDF) + `ForwardDiff.jl` to compute body normal, curvature, and velocity



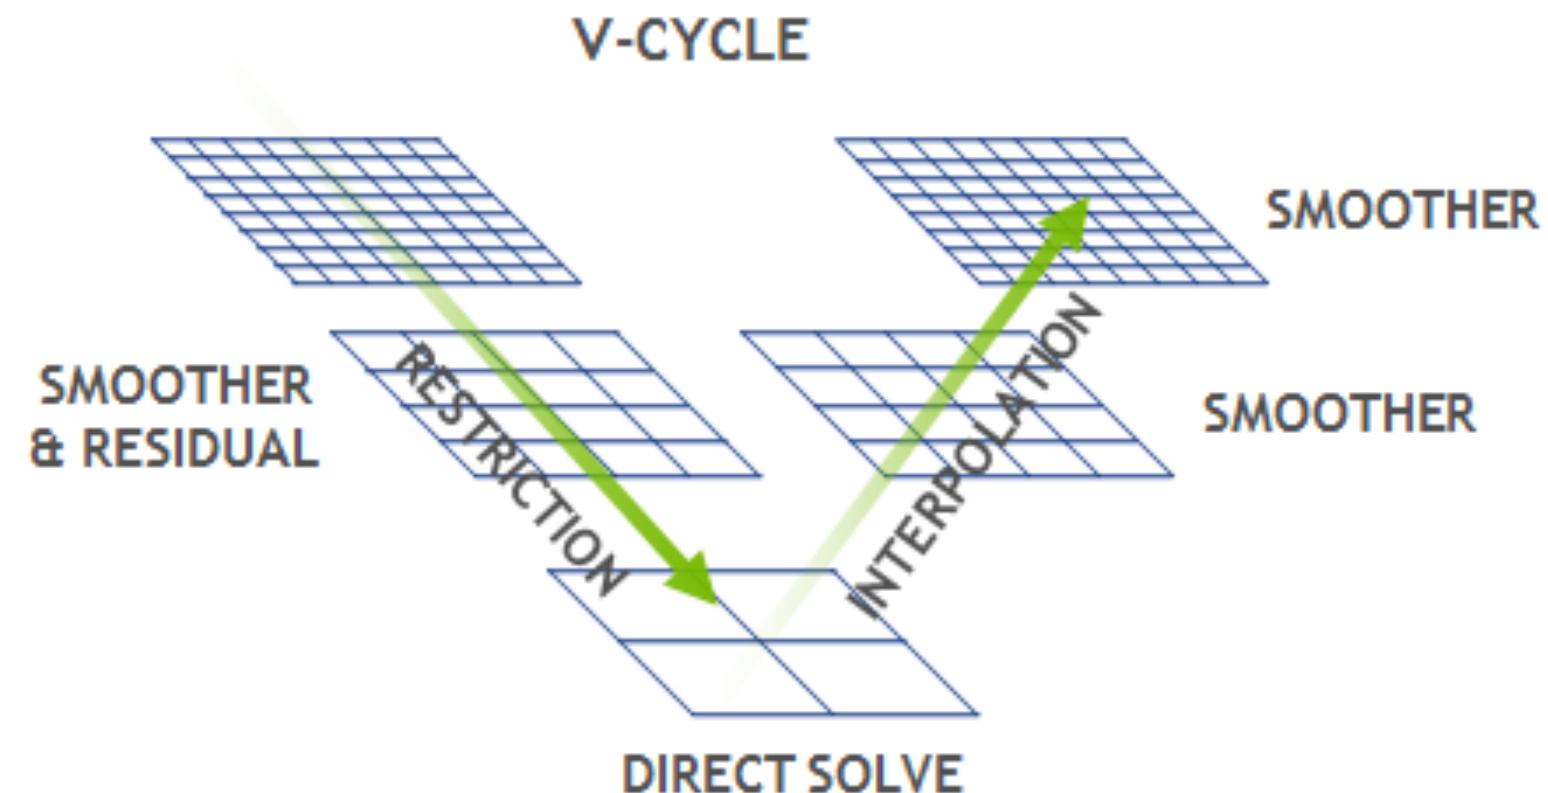
Numerical methods in WaterLily.jl

- Incompressible fluid flow
- Finite volume in a staggered grid
- Boundary data immersion method for solid geometries
 - Signed distance function (SDF) + `ForwardDiff.jl` to compute body normal, curvature, and velocity



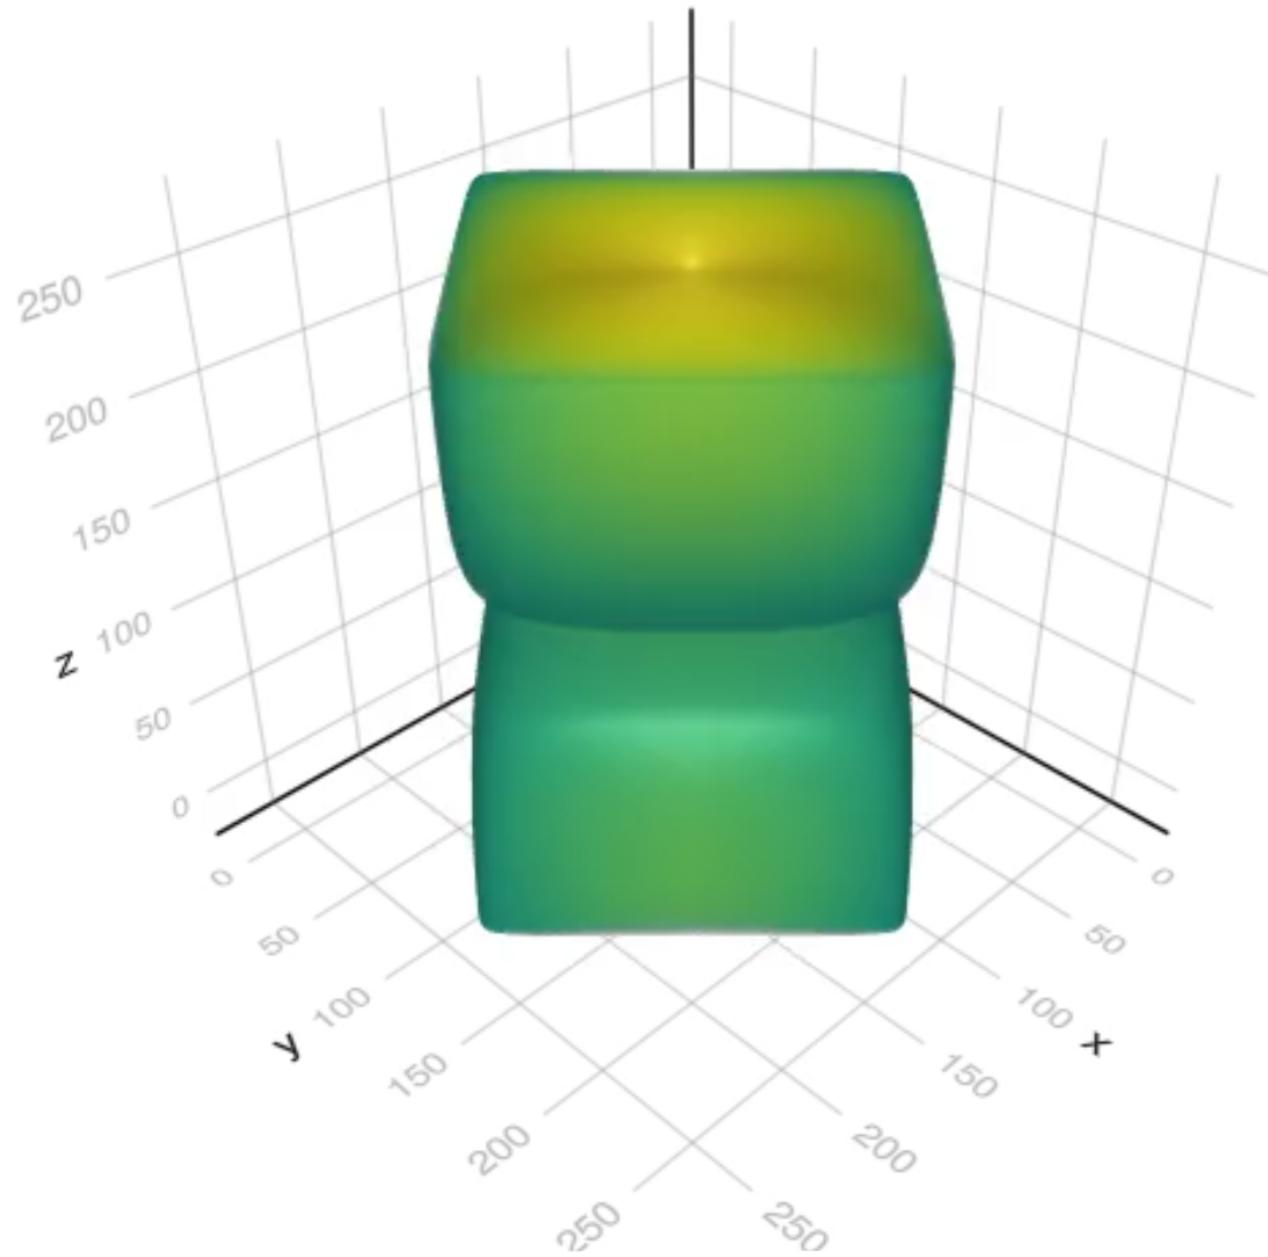
Numerical methods in WaterLily.jl

- Incompressible fluid flow
- Finite volume in a staggered grid
- Boundary data immersion method for solid geometries
 - Signed distance function (SDF) + ForwardDiff.jl to compute body normal, curvature, and velocity
- Geometric multigrid for the pressure solver (Poisson equation) + projection method



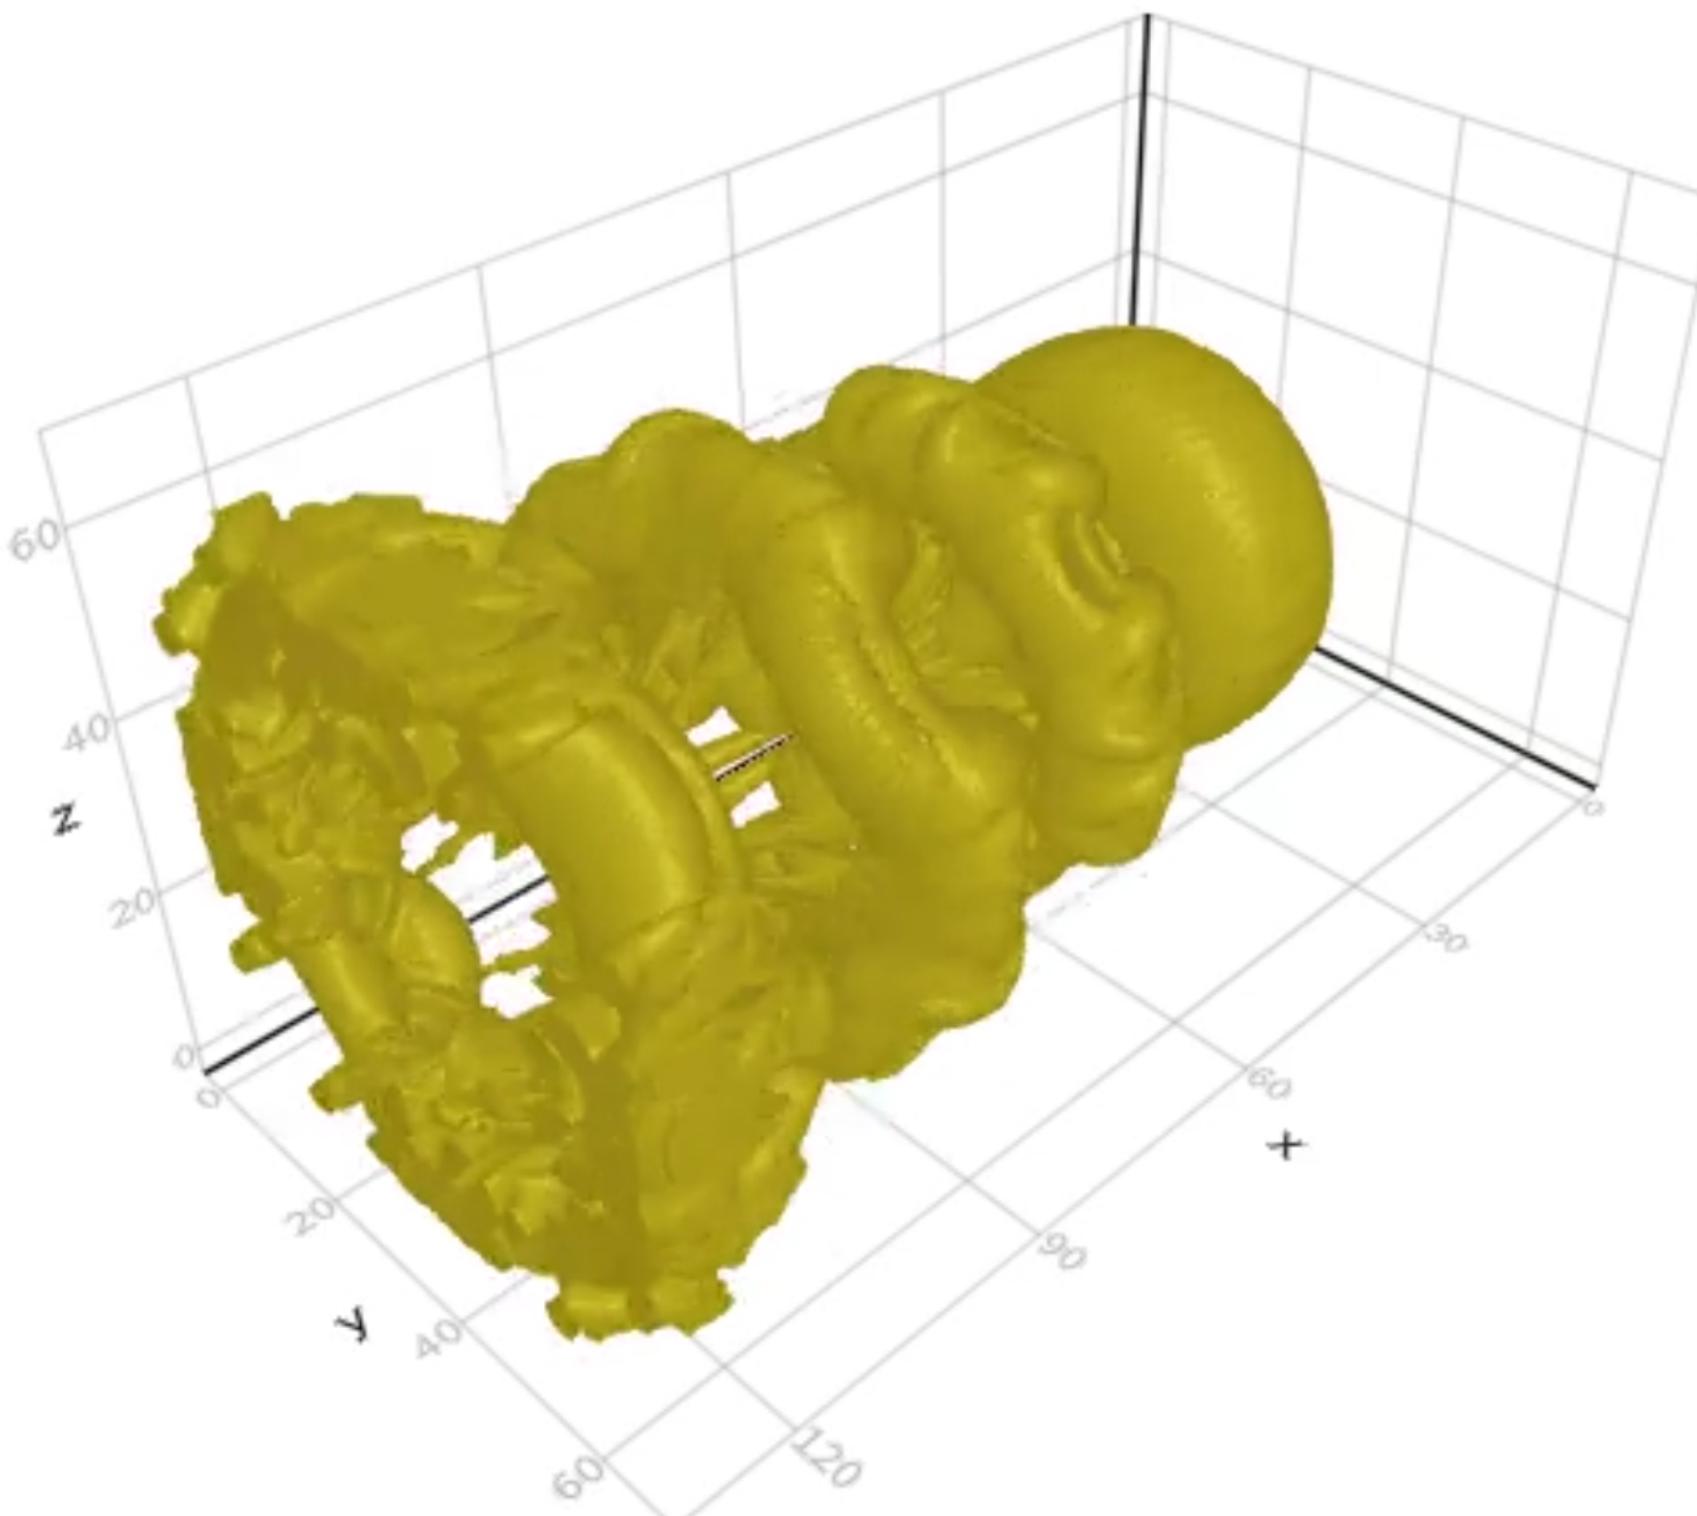
Easy simulations setup

- The `Simulation` struct holds information of both `Fluid` and `Body`
- The user defines the initial condition of the problem, the size of the simulation, and the body SDF (if any).
- The simulation can be run on CPU or GPU just specifying the array type (`Array`, `CuArray`, `ROCArray`, ...)



```
using WaterLily

function TGV(; pow=6, Re=1e5, T=Float32, mem=Array)
    # Define vortex size, velocity, viscosity
    L = 2^pow; U = 1; v = U*L/Re
    # Taylor-Green-Vortex initial velocity field
    function uλ(i,xyz)
        x,y,z = @. (xyz-1.5)*π/L
        i==1 && return -U*sin(x)*cos(y)*cos(z) # u_x
        i==2 && return U*cos(x)*sin(y)*cos(z) # u_y
        return 0.                                # u_z
    end
    # Initialize simulation
    return Simulation((L, L, L), (0, 0, 0), L; U, uλ, v, T, mem)
end
```



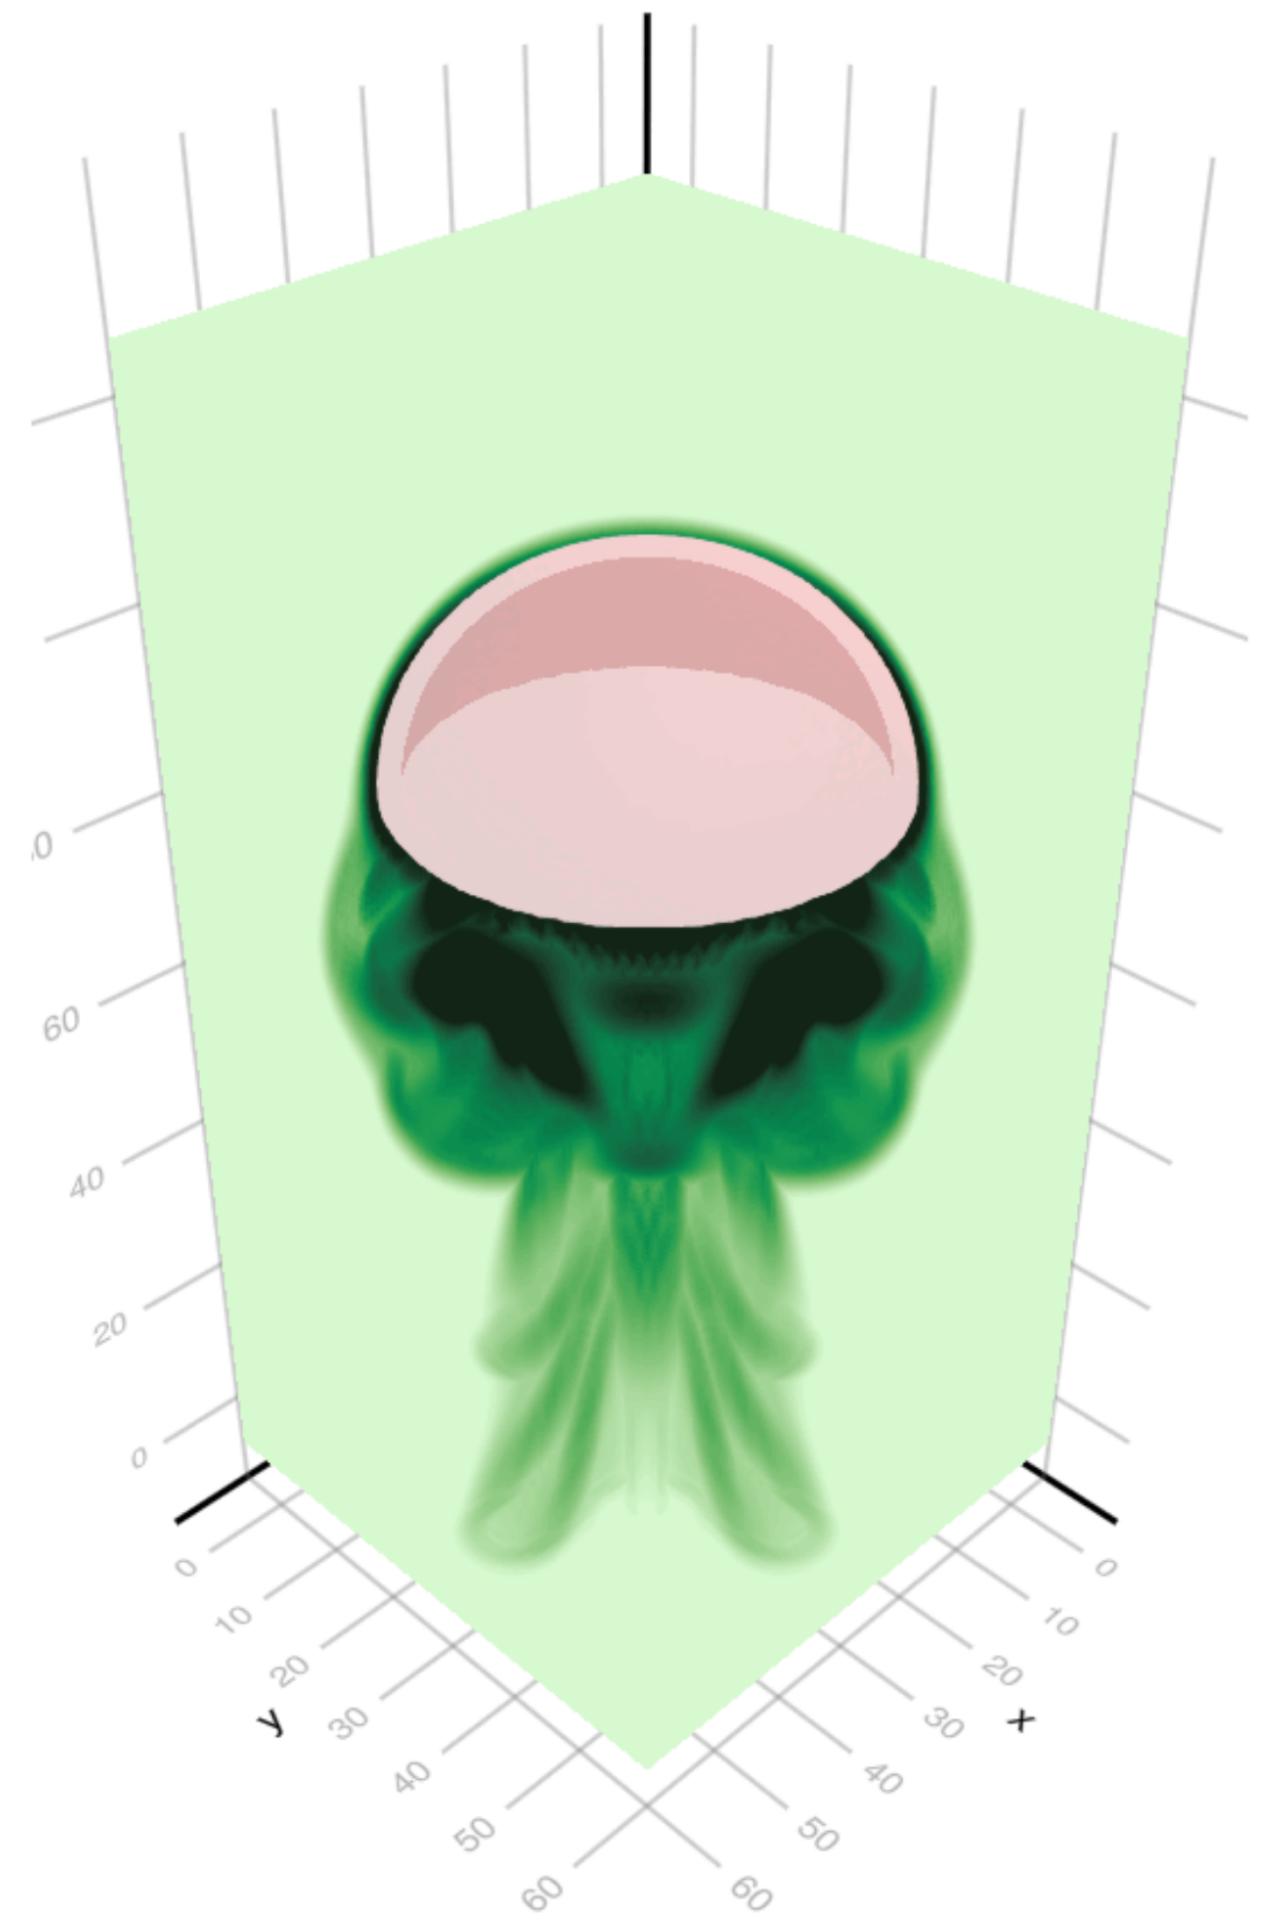
```
using WaterLily
using StaticArrays

function donut(p=6;Re=1e3,mem=Array,U=1)

    # Define simulation size, geometry dimensions, viscosity
    n = 2^p
    center,R,r = SA[n/2,n/2,n/2], n/4, n/16
    v = U*R/Re

    # Apply signed distance function for a torus
    norm2(x) = √sum(abs2,x)
    body = AutoBody() do xyz,t
        x,y,z = xyz - center
        norm2(SA[x,norm2(SA[y,z])-R])-r
    end

    # Initialize simulation and return center for flow viz
    Simulation((2n,n,n),(U,0,0),R;v,mem),center
end
```



```

using WaterLily
using StaticArrays
function jelly(p=5;Re=5e2,mem=Array,U=1)

    # Define simulation size, geometry dimensions, & viscosity
    n = 2^p; R = 2n/3; h = 4n-2R; v = U*R/Re

    # Motion functions
    ω = 2U/R

    @fastmath @inline A(t) = 1 .- SA[1,1,0]*0.1*cos(ω*t)
    @fastmath @inline B(t) = SA[0,0,1]*((cos(ω*t)-1)*R/4-h)
    @fastmath @inline C(t) = SA[0,0,1]*sin(ω*t)*R/4

    # Build jelly from a mapped sphere and plane
    sphere = AutoBody((x,t)->abs(sqrt(sum(abs2,x))-R)-1, # sdf
                      (x,t)->A(t).*x+B(t)+C(t)) # map
    plane = AutoBody((x,t)->x[3]-h,(x,t)->x+C(t))
    body = sphere-plane

    # Return initialized simulation
    Simulation((n,n,4n),(0,0,-U),R;v,mem,T=Float32)

end

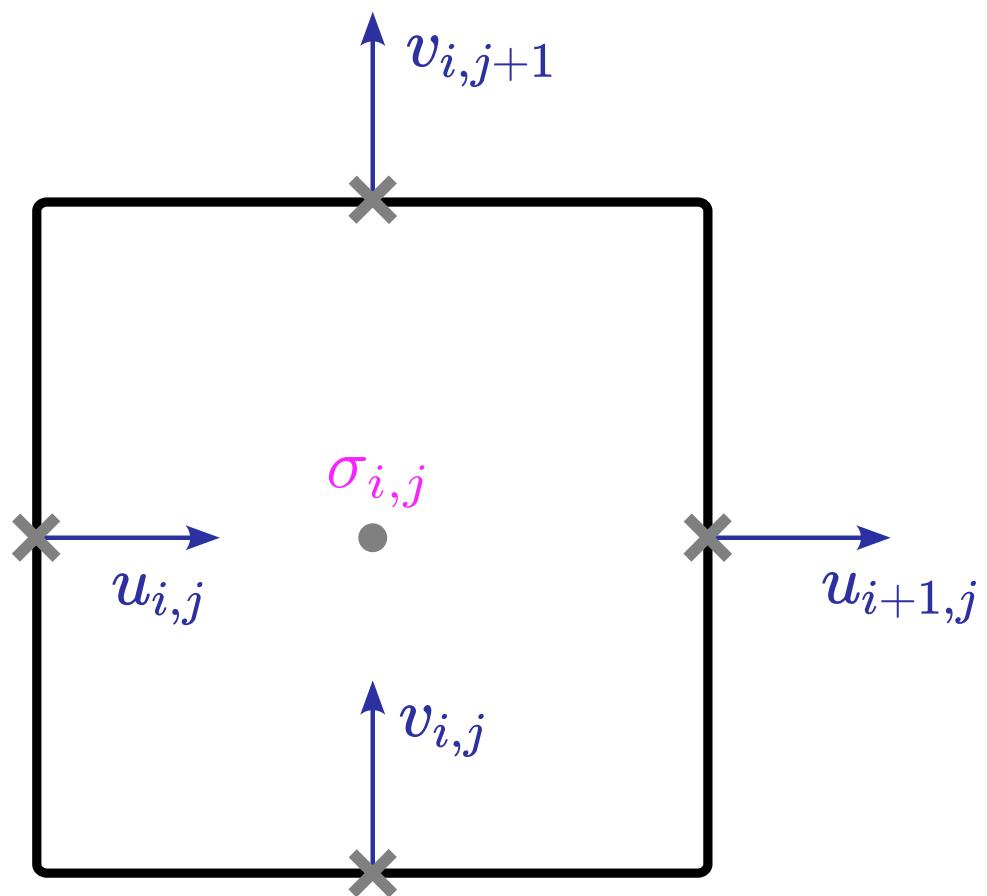
```

Behind the scenes | Loop abstractions

$$\sigma_{i,j} = \iiint \nabla \cdot \mathbf{u} \, dV = \iint \mathbf{u} \cdot \hat{\mathbf{n}} \, dS = (u_{i+1,j} - u_{i,j}) + (v_{i,j+1} - v_{i,j})$$

```
function divergence!(σ, u)
    for d ∈ 1:ndims(σ)
        @loop σ[I] += ∂(d, I, u) over I ∈ inside(σ)
    end
end

# serial loop macro
macro loop(args...)
    ex,_,itr = args # gets expression and iterator info
    op,I,R = itr.args # from iterator info, get index and range
    @assert op ∈ (:(:), :(in))
    return quote
        for $I ∈ $R
            $ex # contains I
        end
    end |> esc
end
```



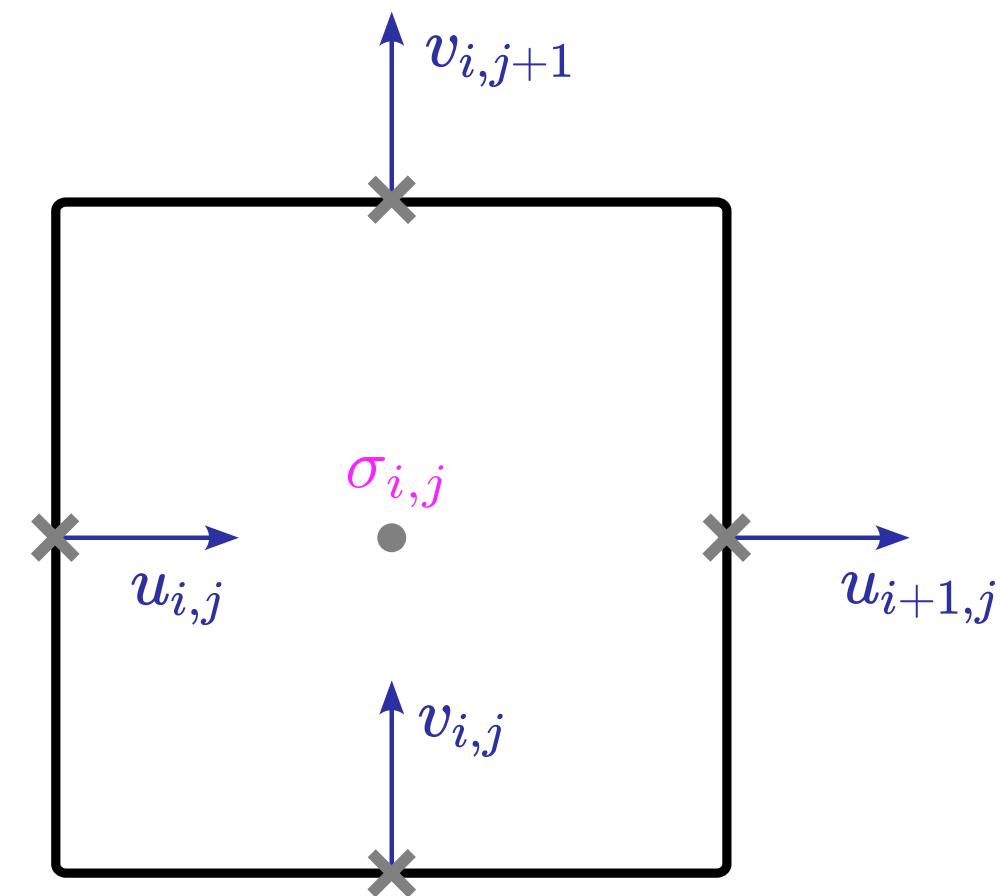
Behind the scenes | Loop abstractions + KernelAbstractions.jl

$$\sigma_{i,j} = \iiint \nabla \cdot \mathbf{u} \, dV = \iint \mathbf{u} \cdot \hat{\mathbf{n}} \, dS = (u_{i+1,j} - u_{i,j}) + (v_{i,j+1} - v_{i,j})$$

```
function divergence!(σ, u)
    for d ∈ 1:ndims(σ)
        @loop σ[I] += ∂(d, I, u) over I ∈ inside(σ)
    end
end
```

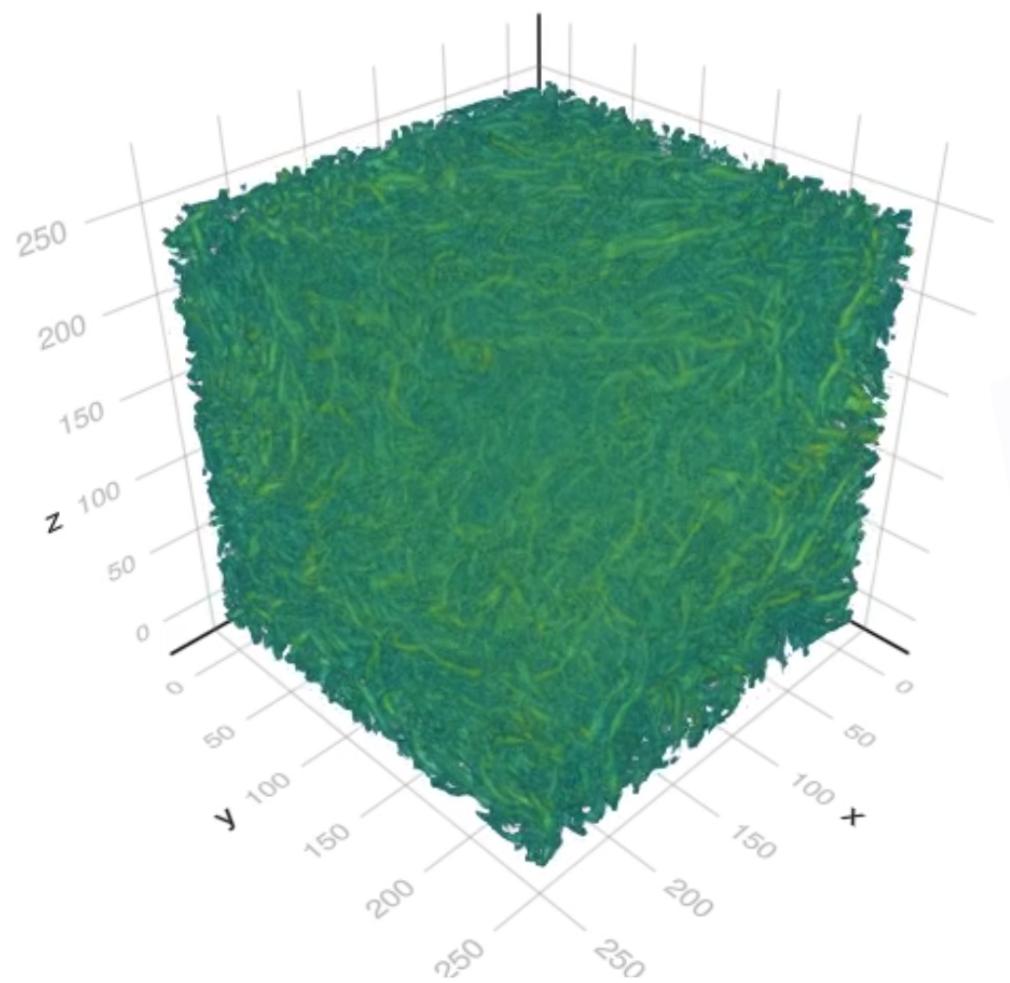
```
# KA-adapted loop macro
macro loop(args...)
    ex, _, itr = args
    _, I, R = itr.args; sym = []
    grab!(sym, ex)      # get arguments and replace composites in `ex`
    setdiff!(sym, [I]) # don't want to pass I as an argument
    @gensym kern       # generate unique kernel function name
    return quote
        @kernel function $kern($rep.(sym)...),@Const(I₀)) # replace composite arguments
            $I = @index(Global, Cartesian)
            $I += I₀ # offset
            $ex # contains I
        end
        $kern(get_backend($(sym[1])), 64)($sym..., $R[1]-oneunit($R[1]), ndrange=size($R))
    end |> esc
end
```

Codebase from 1000 LOC to... 1000LOC!

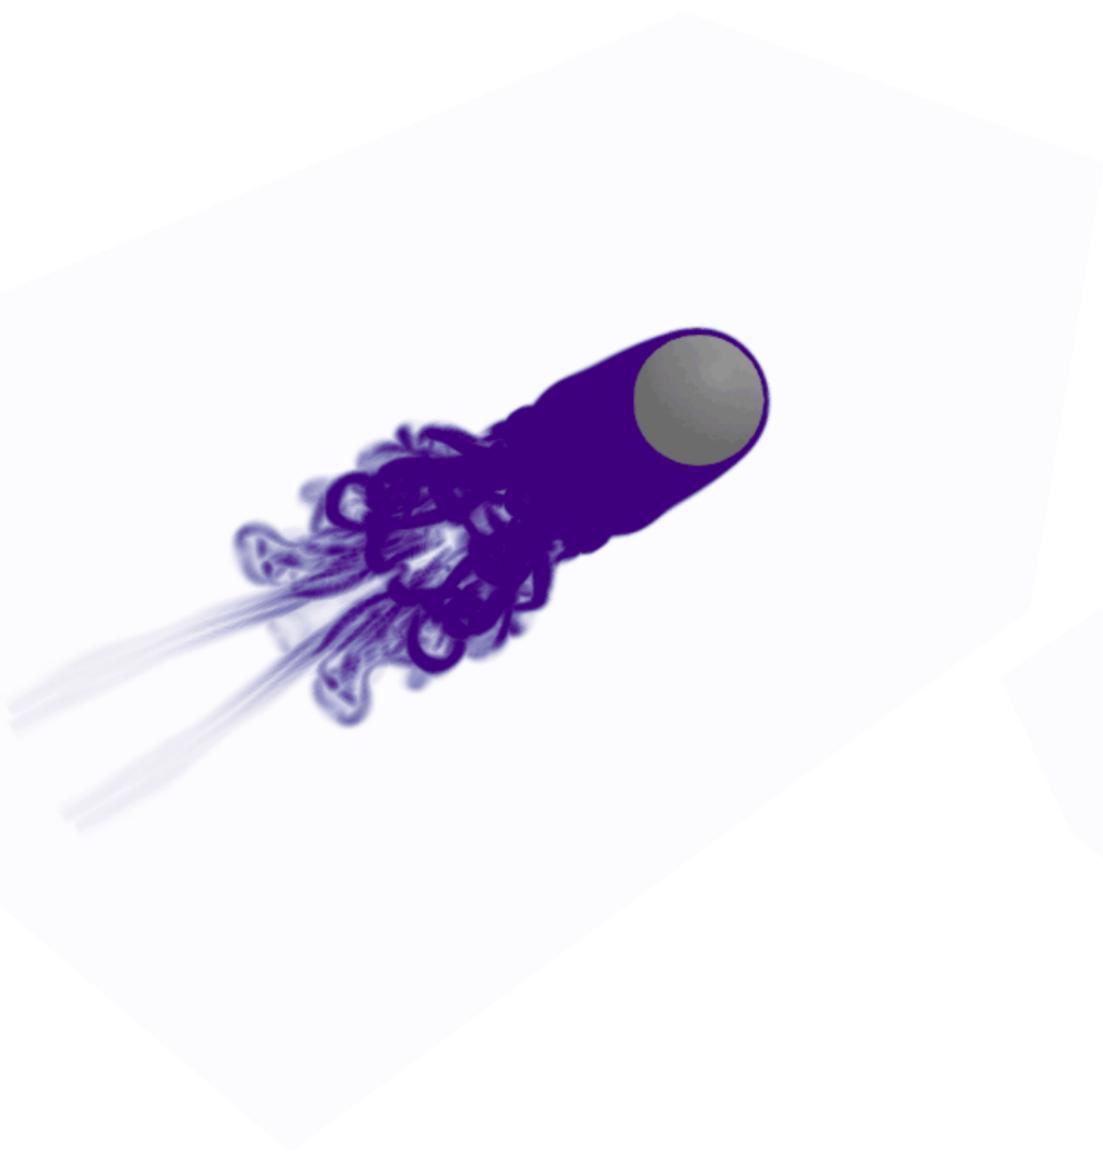


Cases | Validation, Benchmarking & Profiling

TGV

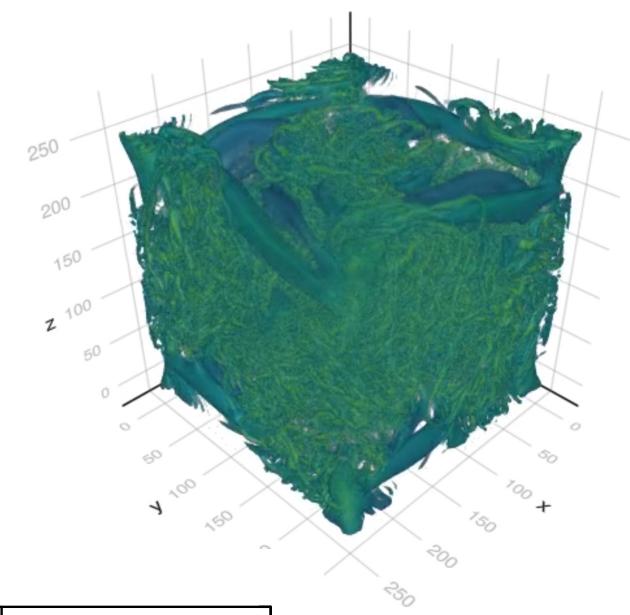


Sphere

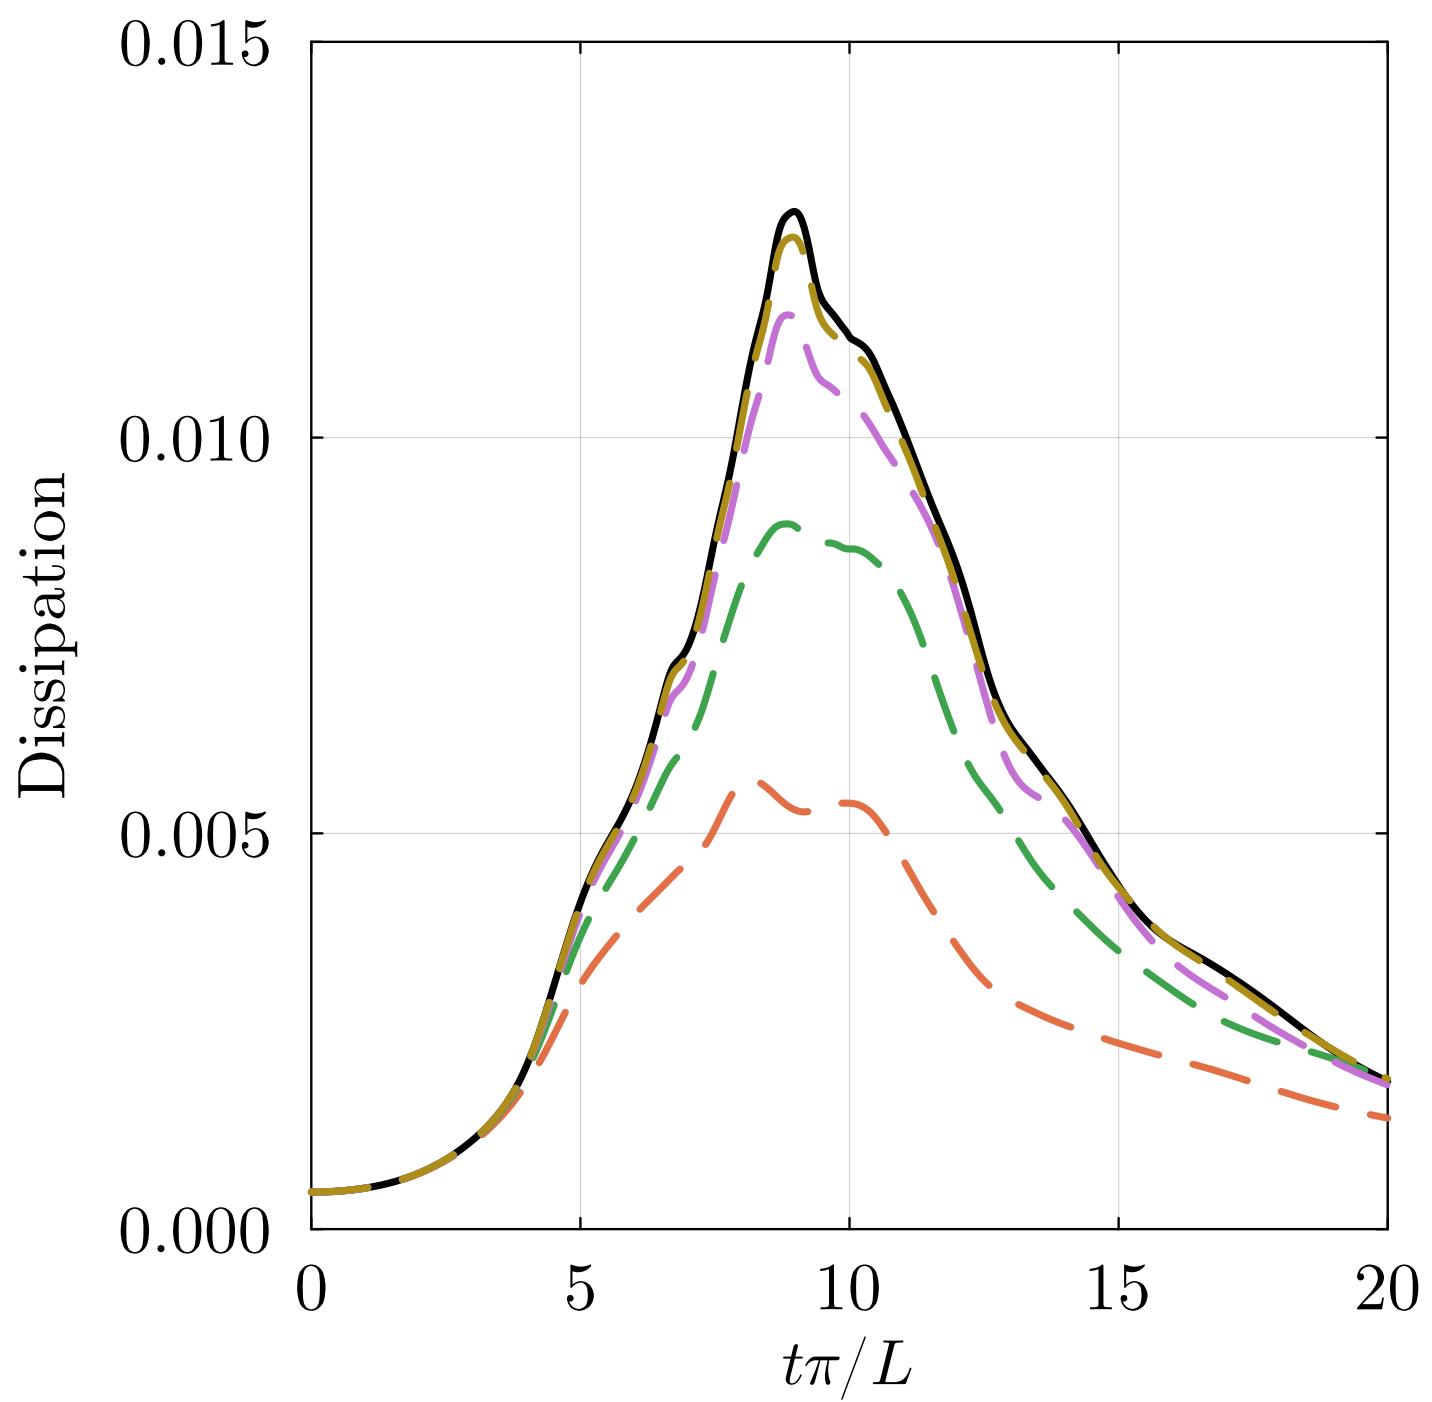
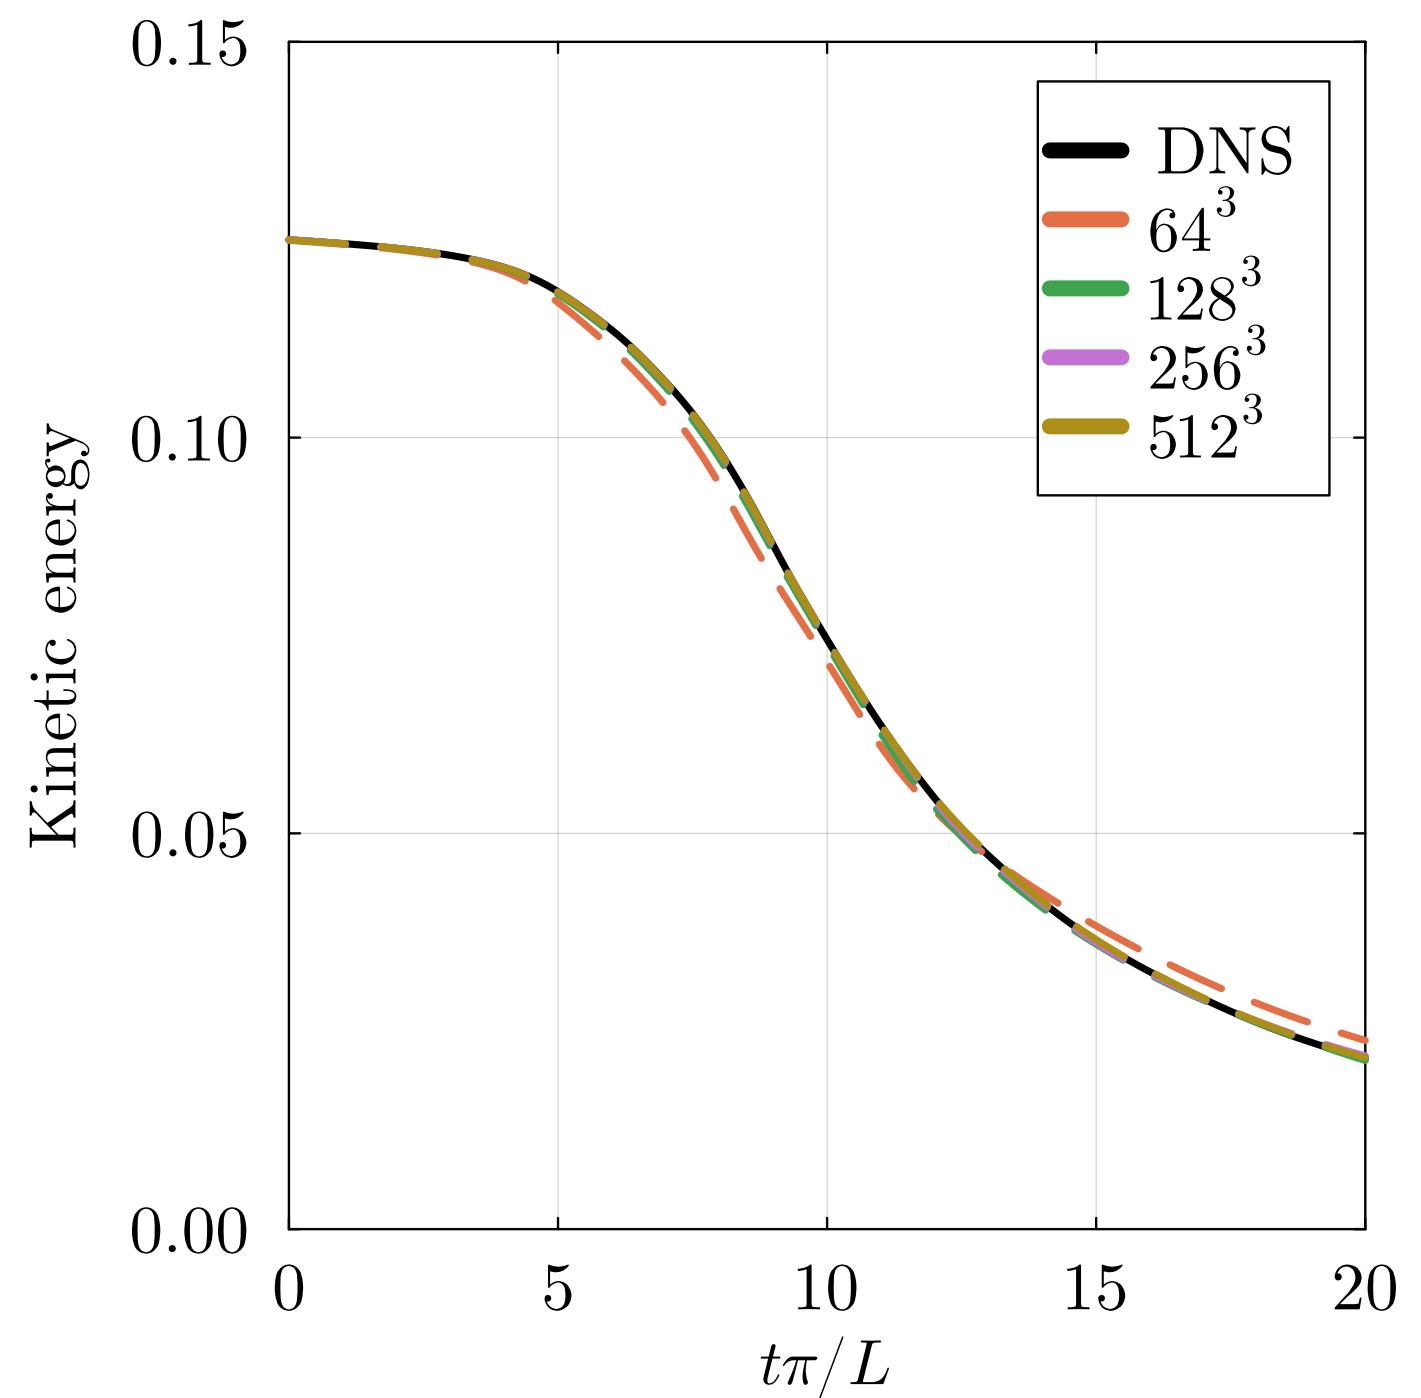


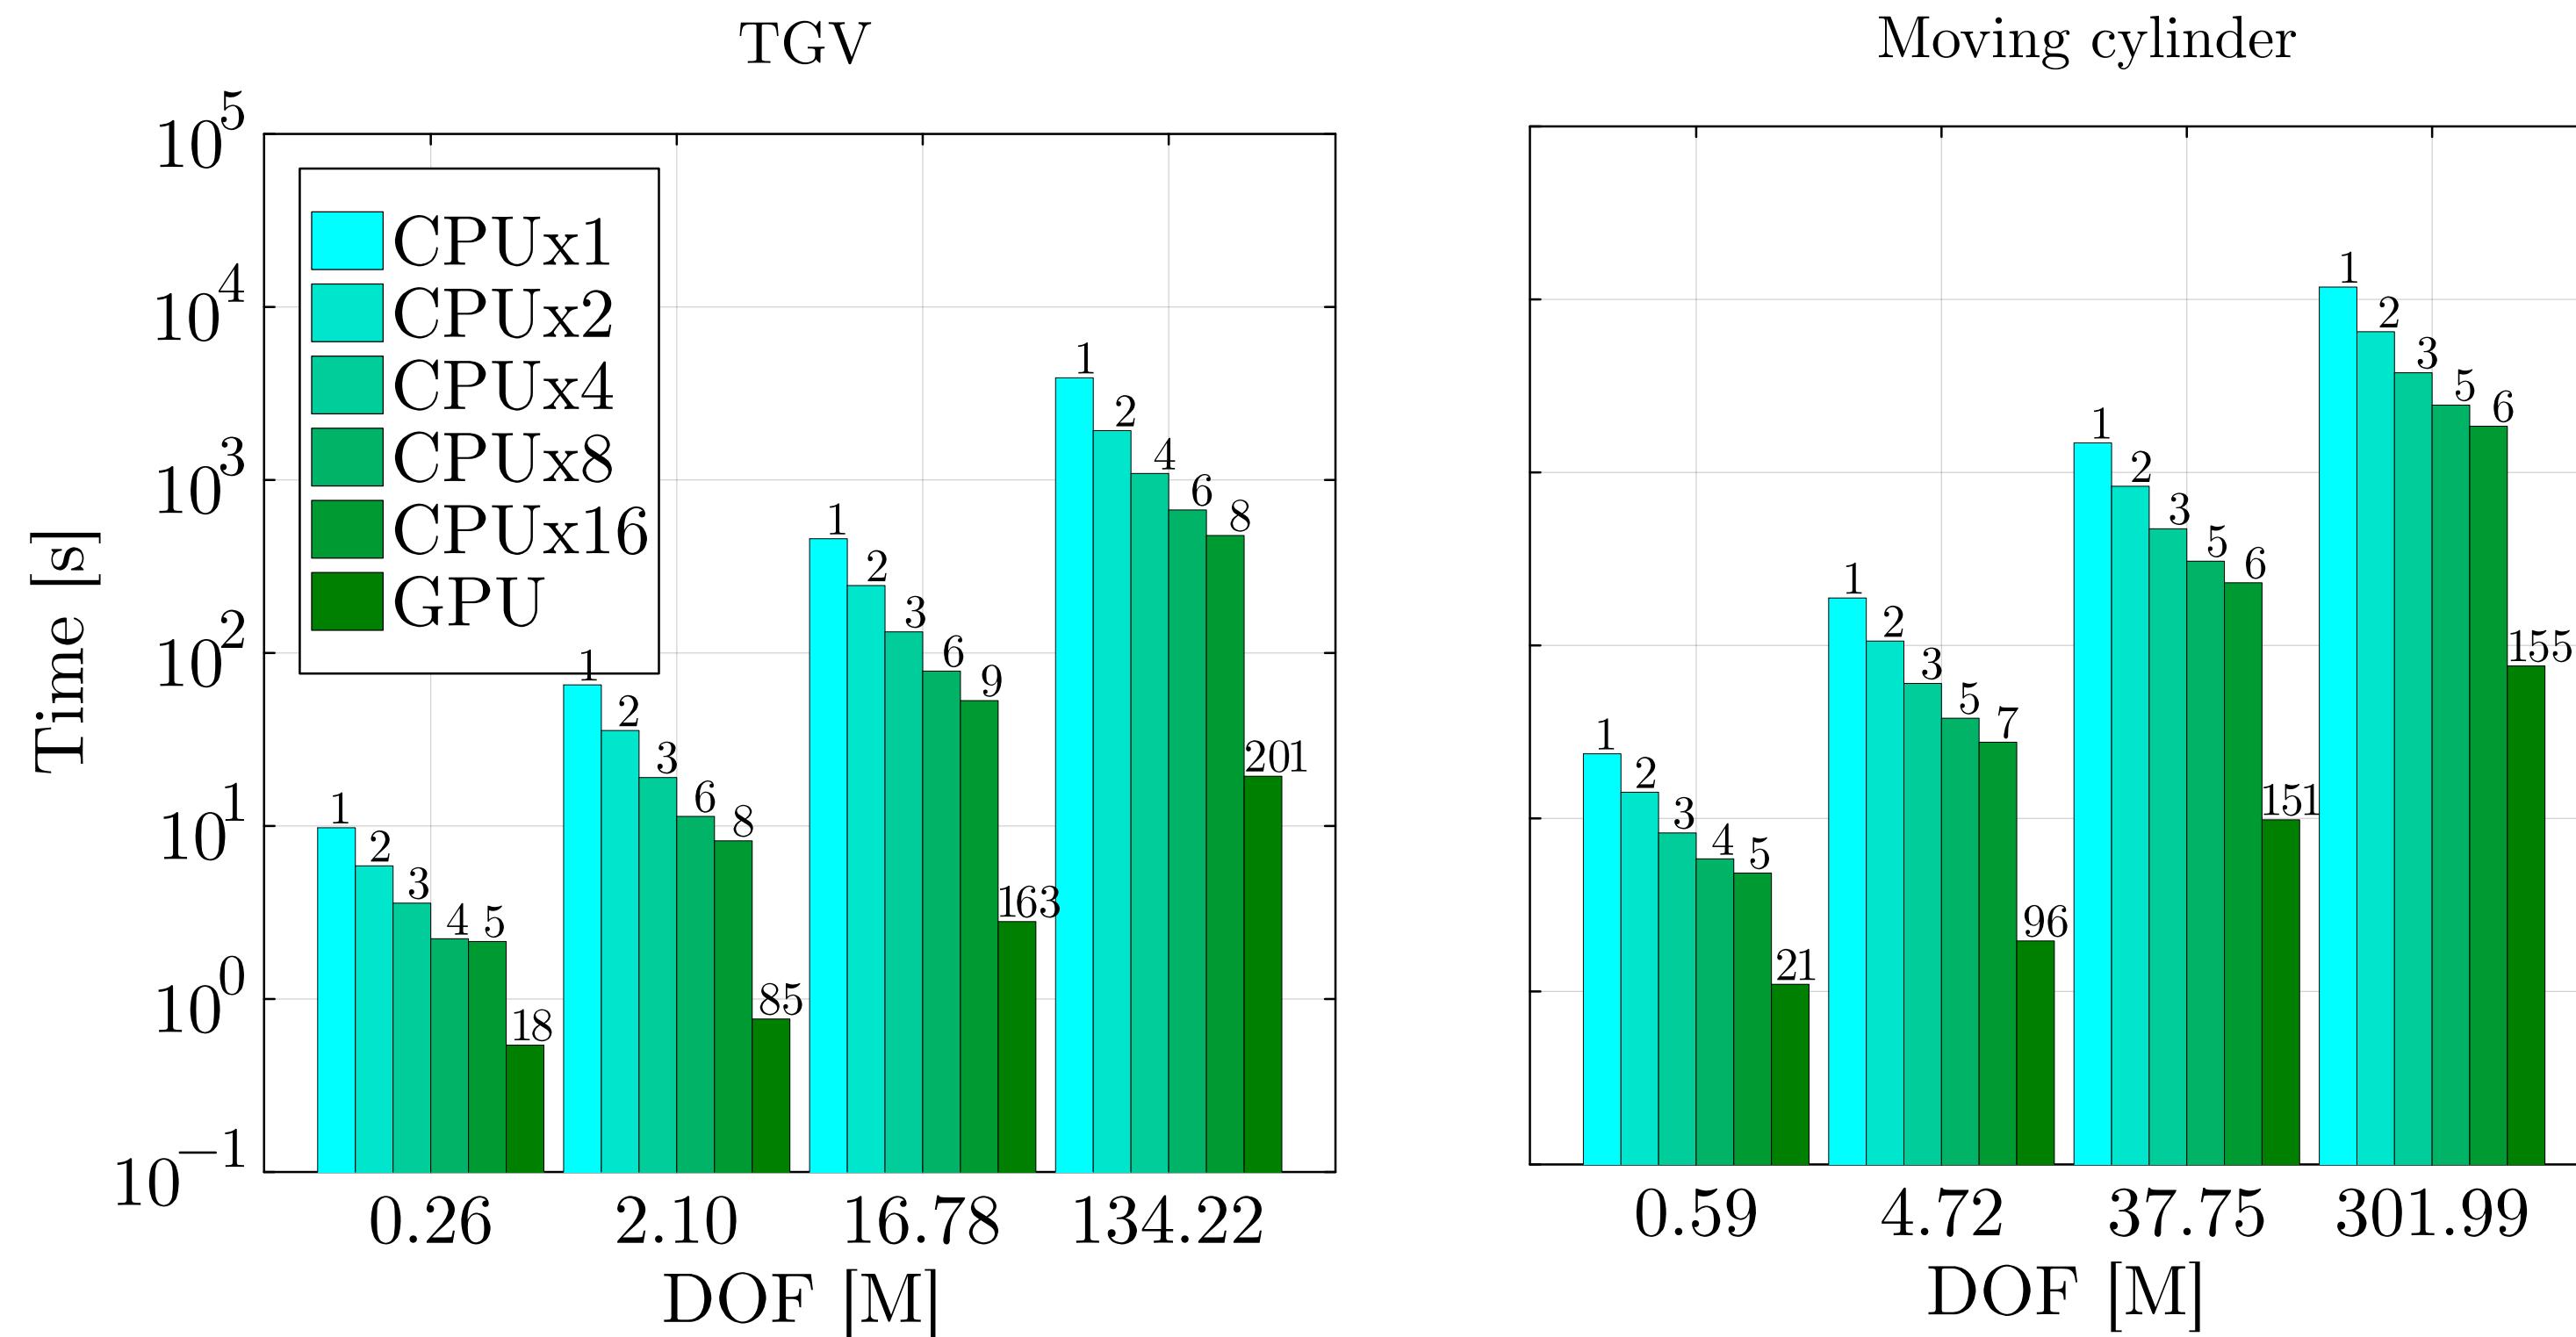
Moving cylinder

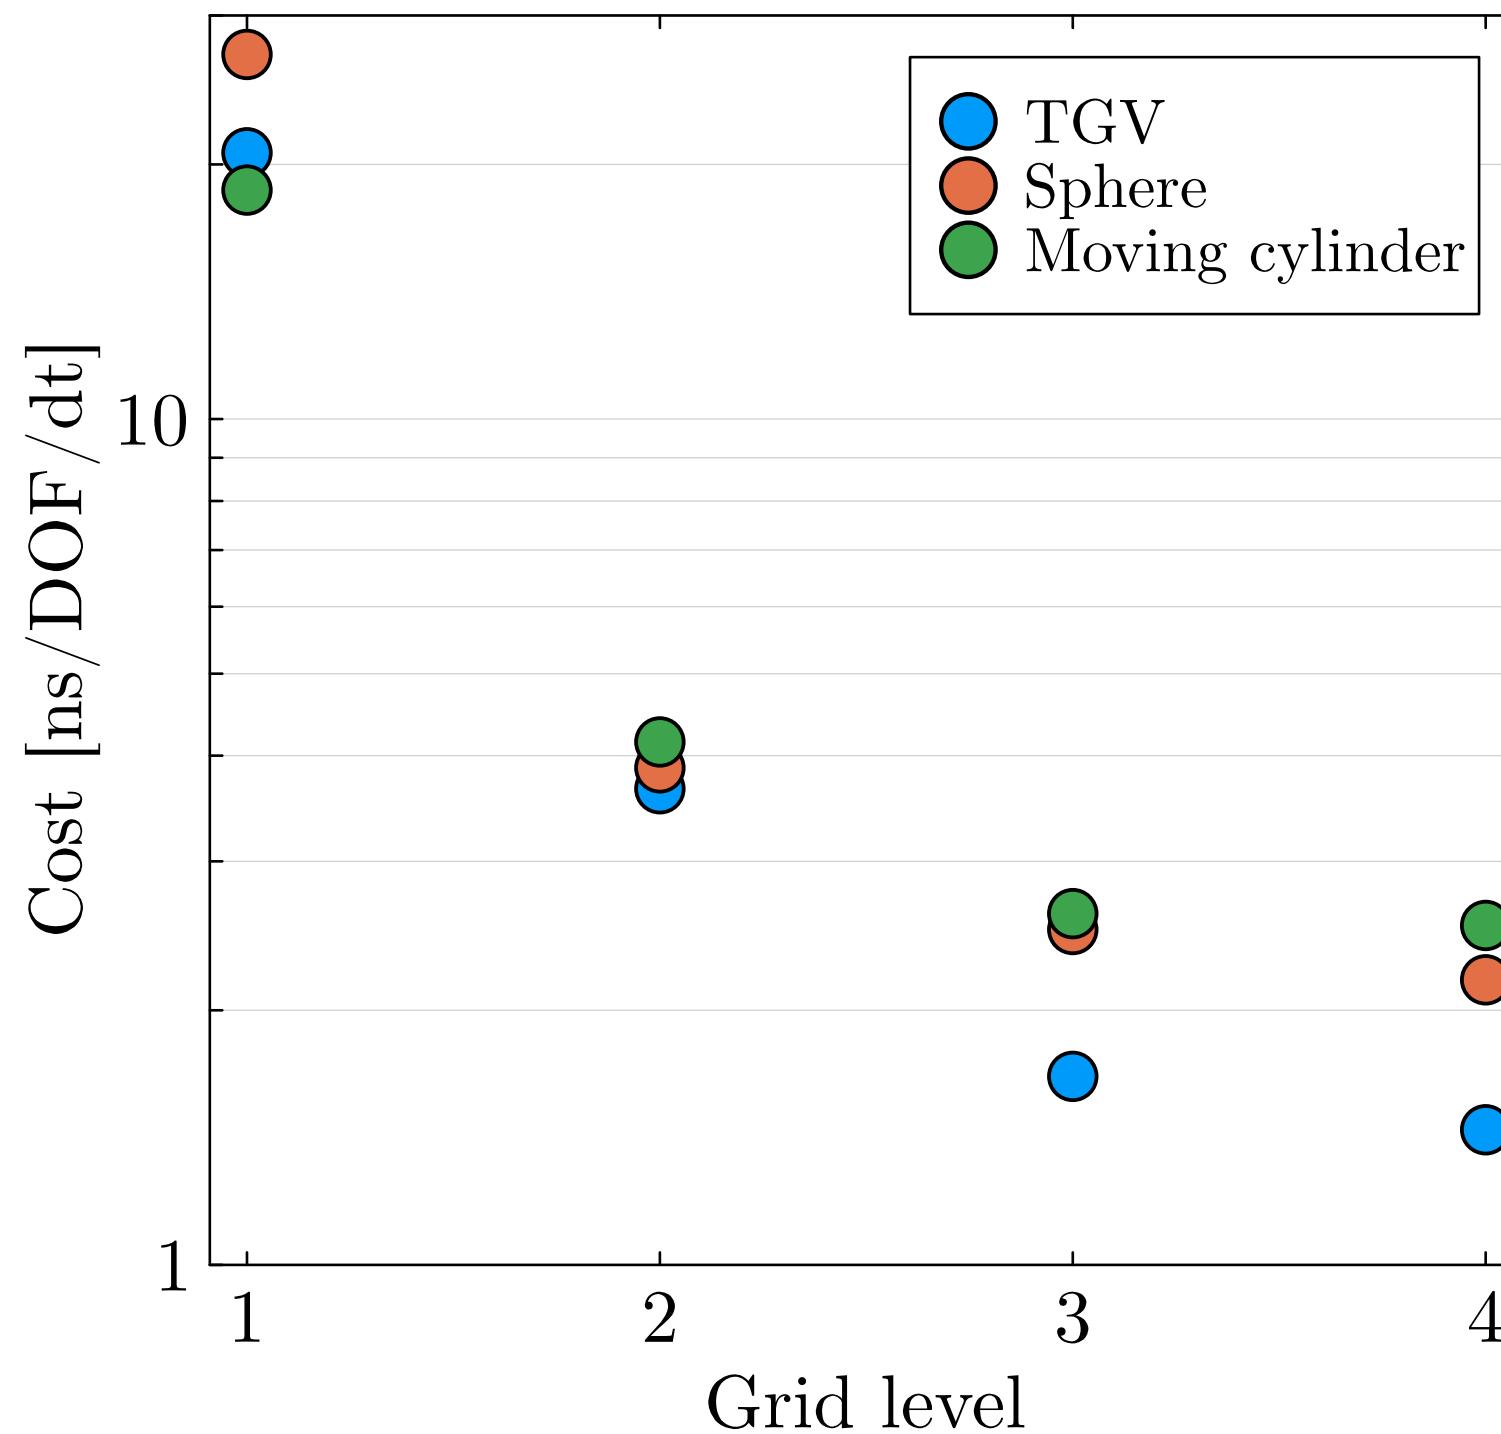


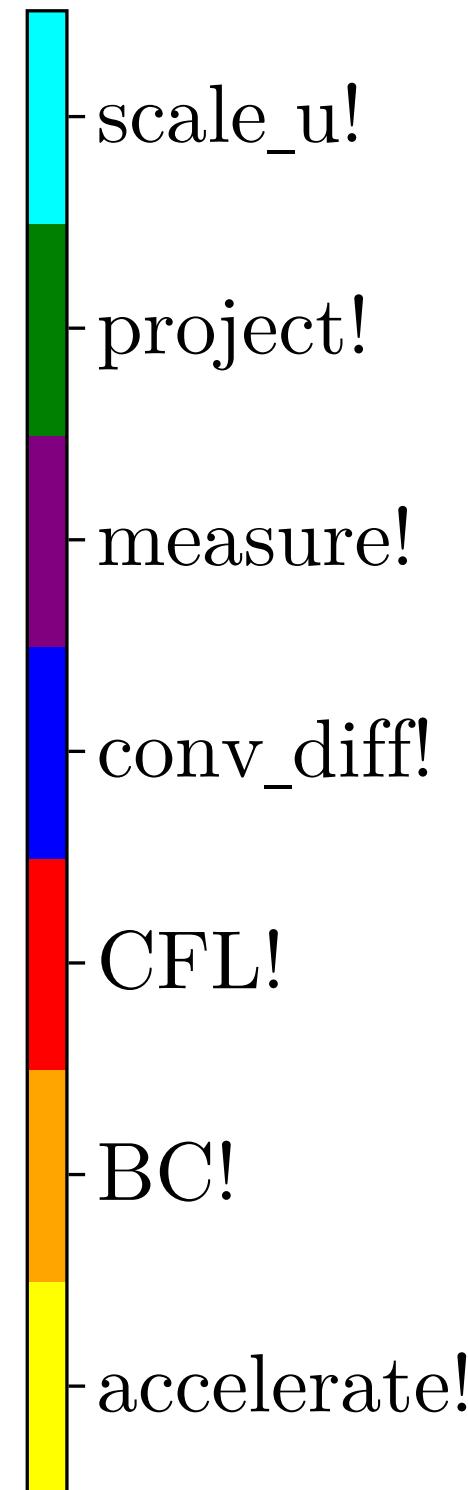
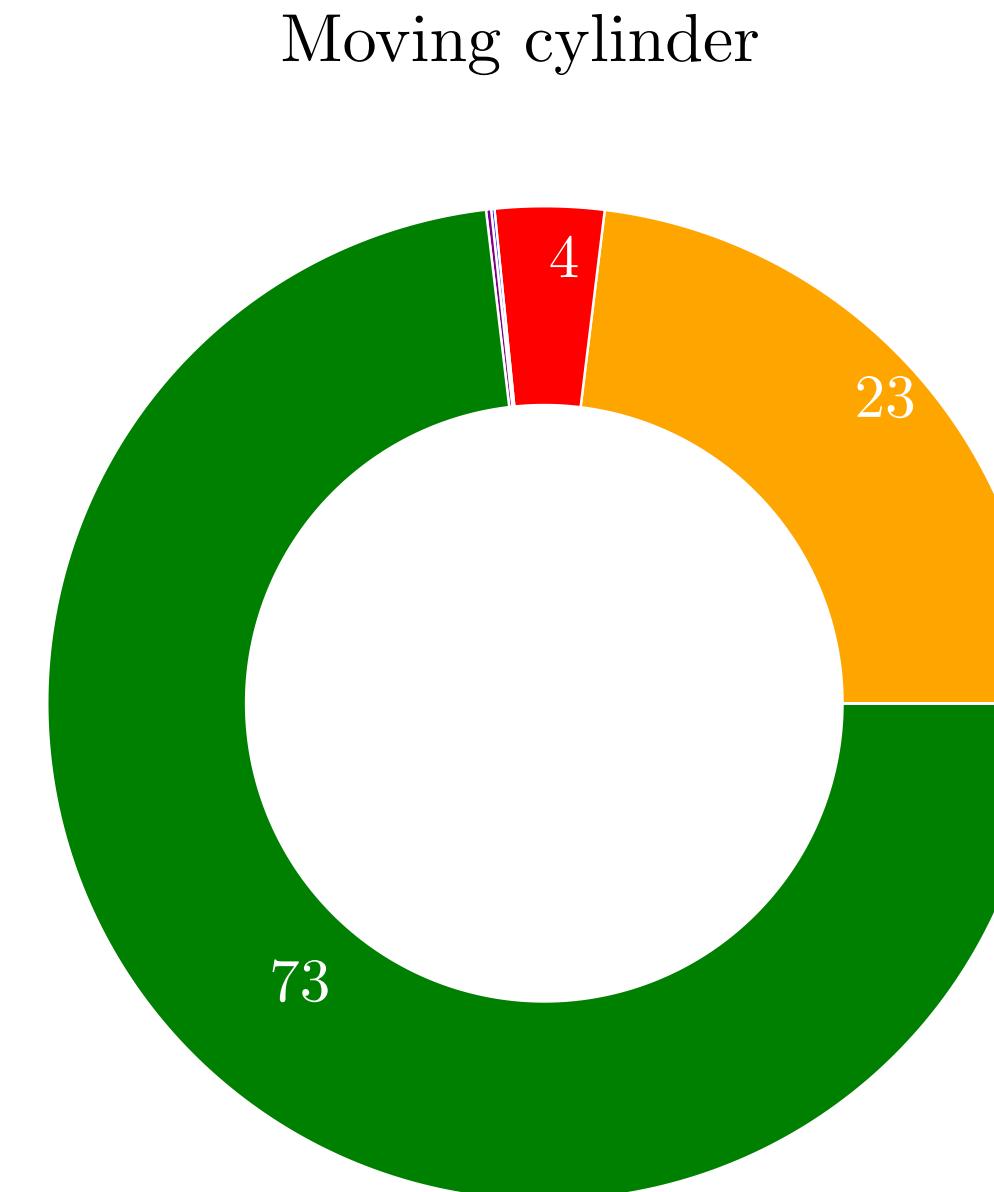
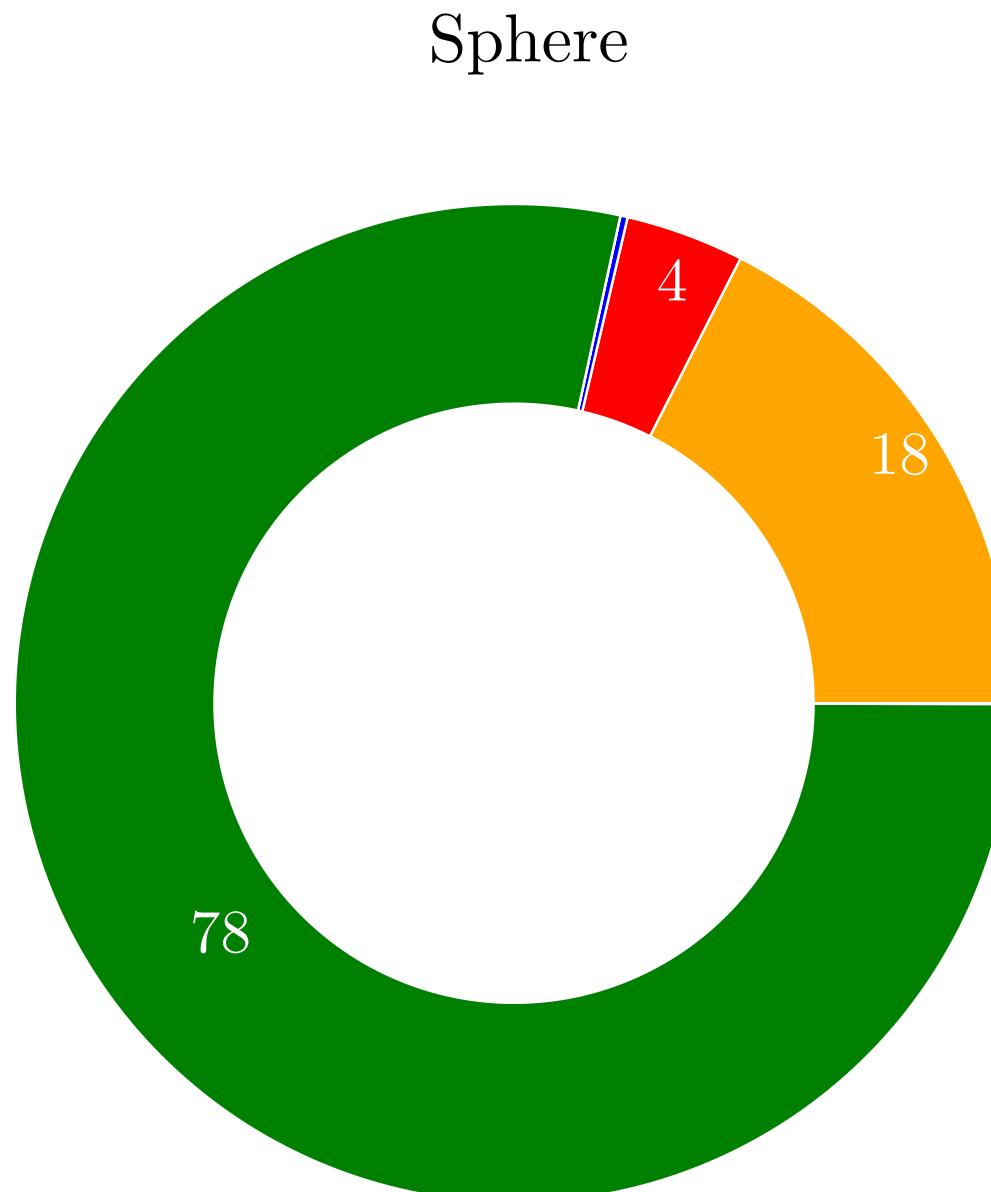
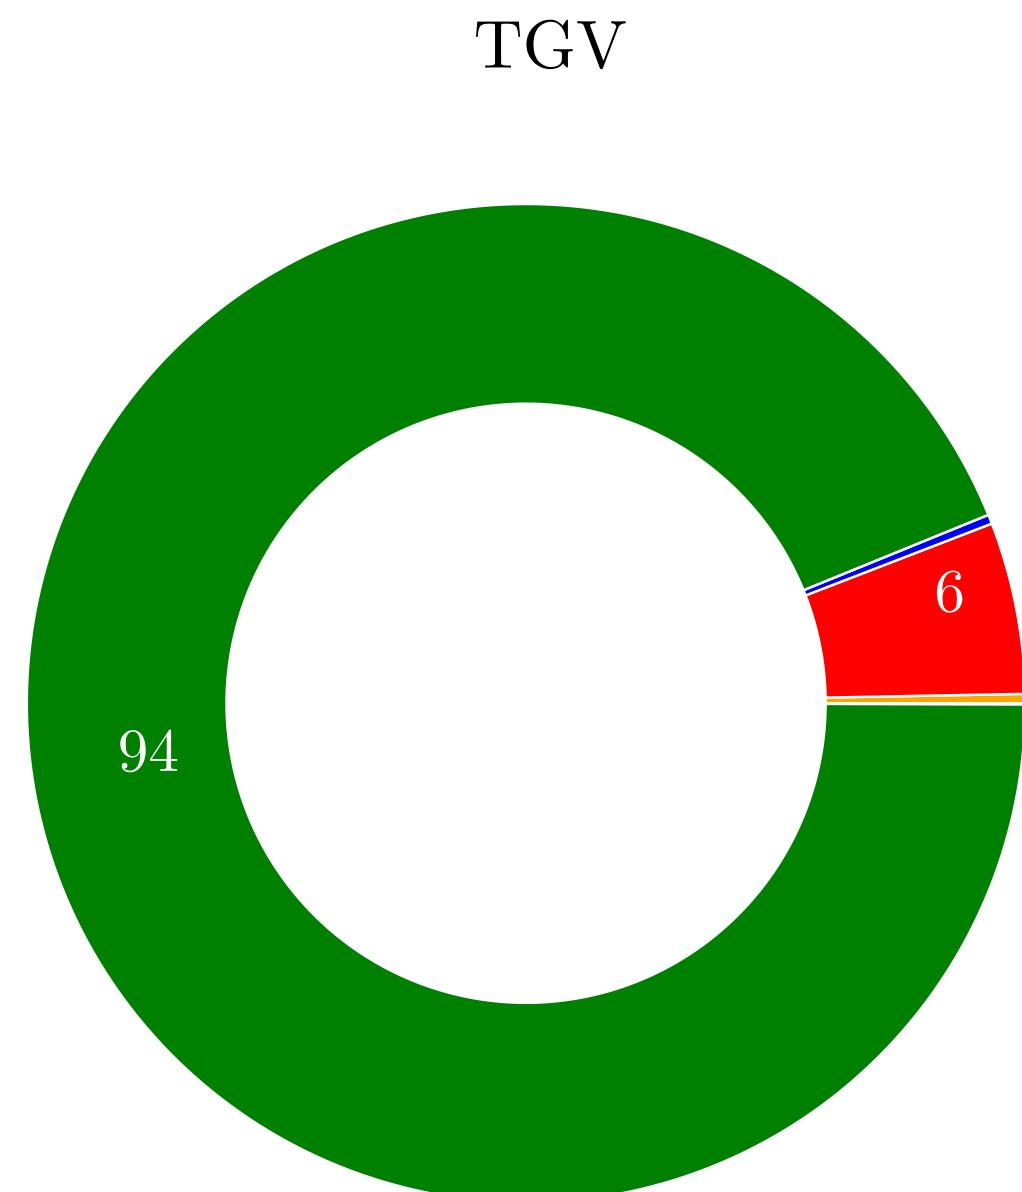


TGV







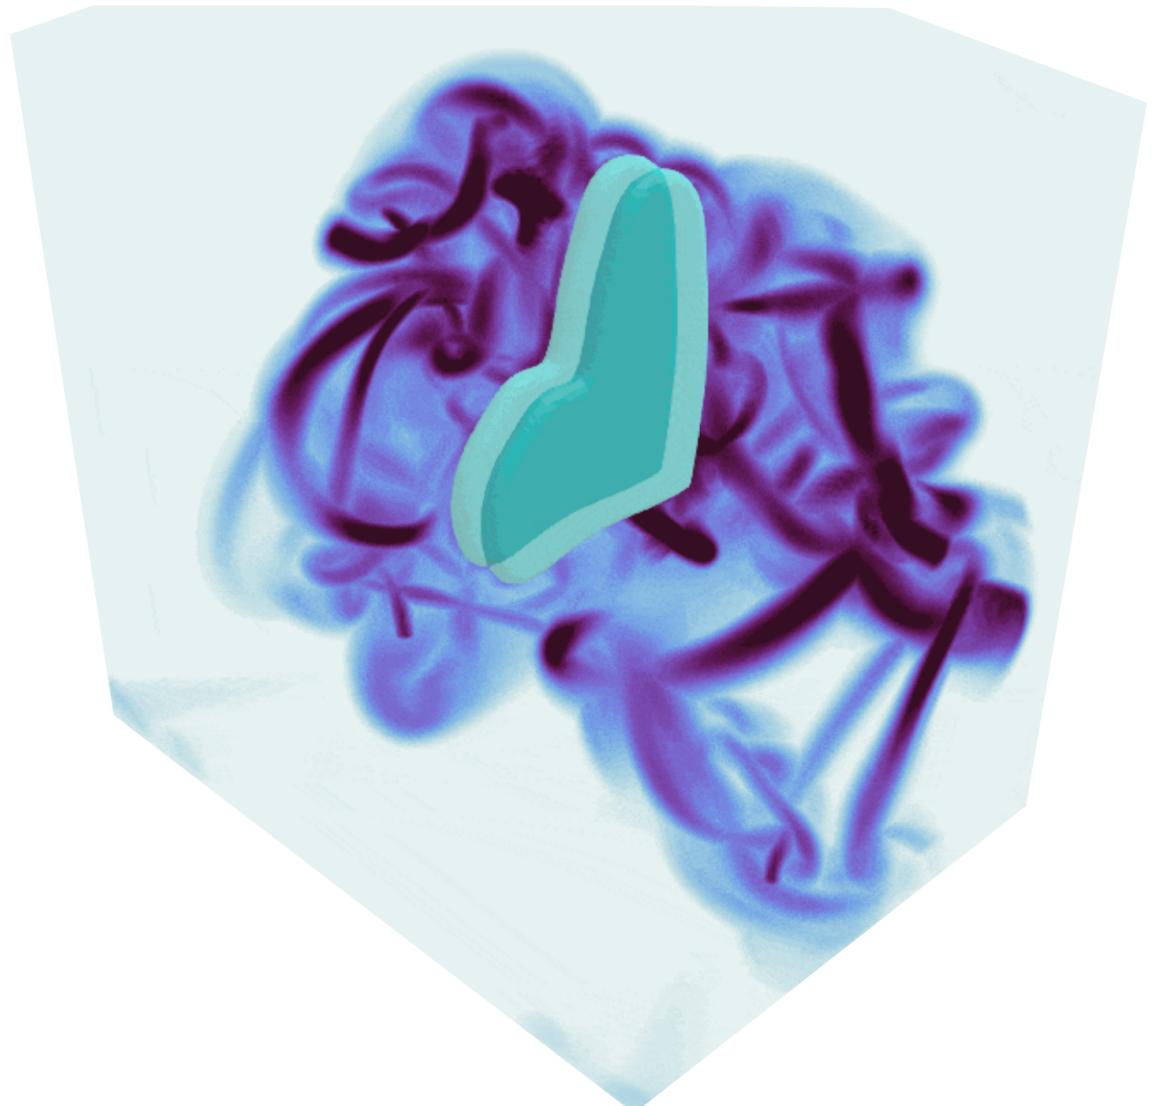


WaterLily.jl recap

- CFD solver with backend-agnostic execution thanks to `KernelAbstractions.jl`
- Simulation of dynamic bodies using immersed boundary method and signed distance function
- Speedups up to 200x on a cluster-grade GPU
- Competitive cost [ns/DOF/dt] wrt well-established CFD solvers written in C/C++/Fortran
- AD-ready on CPU backend using `ForwardDiff.jl`

Coming soon (and not so soon)

- Memory-distributed parallelism (`MPI.jl`)
- AD on both CPU and GPU through `Enzyme.jl`
- Optimized Biot-Savart boundary conditions for confined domains
- Volume-of-fluid (VOF) for multiphase flows
- Grid stretching

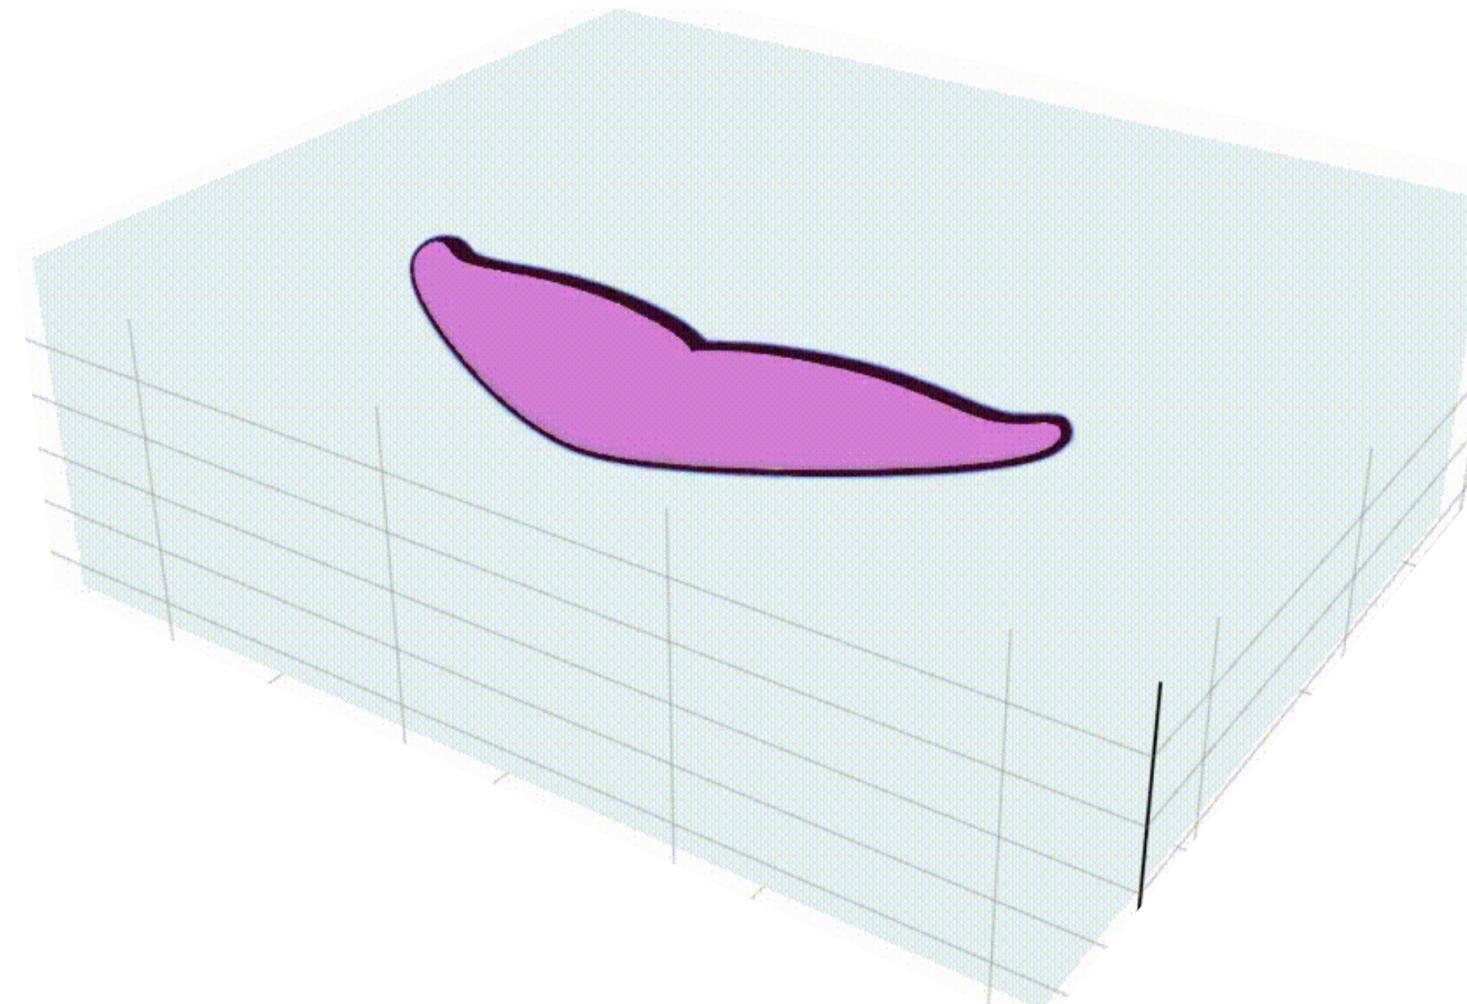


WaterLily.jl: A differentiable solver to simulate fluid flow and dynamic bodies with heterogeneous execution

Bernat Font*, Gabriel D. Weymouth

JuliaCon 2024

WaterLily.jl



Slides

